

## Mastering Spring 5.0

Master reactive programming, microservices, Cloud Native applications, and more

*This book is based on Spring Version 5.0 RC1*

Ranga Rao Karanam

Packt

BIRMINGHAM - MUMBAI

Copyrighted material

## Mastering Spring 5.0

Copyright © 2017 Packt Publishing

All rights reserved. No part of this book may be reproduced, stored in a retrieval system, or transmitted in any form or by any means, without the prior written permission of the publisher, except in the case of brief quotations embedded in critical articles or reviews.

Every effort has been made in the preparation of this book to ensure the accuracy of the information presented. However, the information contained in this book is sold without warranty, either express or implied. Neither the author, nor Packt Publishing, and its dealers and distributors will be held liable for any damages caused or alleged to be caused directly or indirectly by this book.

Packt Publishing has endeavored to provide trademark information about all of the companies and products mentioned in this book by the appropriate use of capitals. However, Packt Publishing cannot guarantee the accuracy of this information.

First published: June 2017

Production reference: 1240617

Published by Packt Publishing Ltd.

Livery Place  
35 Livery Street  
Birmingham  
B3 2PB, UK.

ISBN 978-1-78712-317-5

[www.packtpub.com](http://www.packtpub.com)

Copyrighted material

Some pages are omitted from this book preview.

## www.PacktPub.com

For support files and downloads related to your book, please visit [www.PacktPub.com](http://www.PacktPub.com).

Did you know that Packt offers eBook versions of every book published, with PDF and ePUB files available? You can upgrade to the eBook version at [www.PacktPub.com](http://www.PacktPub.com) and as a print book customer, you are entitled to a discount on the eBook copy. Get in touch with us at [service@packtpub.com](mailto:service@packtpub.com) for more details.

At [www.PacktPub.com](http://www.PacktPub.com), you can also read a collection of free technical articles, sign up for a range of free newsletters and receive exclusive discounts and offers on Packt books and eBooks.



<https://www.packtpub.com/mapt>

Get the most in-demand software skills with Mapt. Mapt gives you full access to all Packt books and video courses, as well as industry-leading tools to help you plan your personal development and advance your career.

### Why subscribe?

- Fully searchable across every book published by Packt
- Copy and paste, print, and bookmark content
- On demand and accessible via a web browser

Copyrighted material

Some pages are omitted from this book preview.

## Table of Contents

<b>Preface</b>	1
<b>Chapter 1: Evolution to Spring Framework 5.0</b>	7
<b>Spring Framework</b>	8
Problems with EJB	8
<b>Why is Spring Framework popular?</b>	8
Simplified unit testing	9
Reduction in plumbing code	9
How does Spring Framework do this magic?	10
Architectural flexibility	11
Keep up with changing times	12
<b>Spring modules</b>	12
Spring Core Container	13
Cross-cutting concerns	13
Web	14
Business	14
Data	15
<b>Spring Projects</b>	15
Spring Boot	16
Spring Cloud	16
Spring Data	16
Spring Batch	17
Spring Security	18
Spring HATEOAS	18
<b>New features in Spring Framework 5.0</b>	19
Baseline upgrades	19
JDK 9 runtime compatibility	20
Usage of JDK 8 features in Spring Framework code	20
Reactive programming support	21

Functional web framework	21
Java modularity with Jigsaw	22
Kotlin support	23
Dropped features	24
<b>Spring Boot 2.0 new features</b>	<b>25</b>
Summary	26

Copyrighted material

<b>Chapter 2: Dependency Injection</b>	<b>27</b>
<b>Understanding dependency injection</b>	<b>28</b>
Understanding dependencies	28
The Spring IoC container	31
Defining beans and wiring	31
Creating a Spring IoC container	33
Java configuration for the application context	33
A quick review	34
Launching the application context with Java configuration	35
The complete configuration	36
The XML configuration for the application context	37
Defining the XML Spring configuration	37
Launching an application context with the XML configuration	38
Writing JUnit using the Spring context	38
Unit testing with mocks	40
Container managed beans	42
Dependency injection types	43
The setter injection	43
The constructor injection	44
Constructor versus setter injection	44
Spring bean scopes	44
Java versus XML configuration	45
The @Autowired annotation in depth	46
The @Primary annotation	46
The @Qualifier annotation	47
Other important Spring annotations	47
Exploring Contexts and dependency injection	48
An example of CDI	49
Summary	50
<b>Chapter 3: Building a Web Application with Spring MVC</b>	<b>51</b>

<b>Java web application architecture</b>	<b>52</b>
Model 1 architecture	52
Model 2 architecture	53
Model 2 Front Controller architecture	54
<b>Basic flows</b>	<b>55</b>
Basic setup	55
Adding dependency for Spring MVC	55
Adding DispatcherServlet to web.xml	56
Creating Spring context	57
Flow 1 - Simple controller flow without View	57
Creating a Spring MVC controller	58
Running the web application	58

————— [ ii ] —————

Copyrighted material

Unit testing	59
Setting up the Controller to test	59
Writing the Test method	60
Flow 2 - Simple controller flow with a View	61
Spring MVC controller	61
Creating a View - a JSP	62
View resolver	62
Unit testing	63
Setting up the Controller to test	63
Writing the Test method	64
Flow 3 - Controller redirecting to a View with Model	64
Spring MVC controller	65
Creating a View	65
Unit testing	66
Setting up the Controller to test	66
Writing the Test method	66
Flow 4 - Controller redirecting to a View with ModelAndView	67
Spring MVC controller	67
Creating a View	68
Unit testing	68
Flow 5 - Controller redirecting to a View with a form	69
Creating a command or form backing object	69
The Controller method to show the form	69
Creating the View with a form	70
Controller get method to handle form submit	72
Unit testing	73
Flow 6 - Adding validation to the previous flow	73
Hibernate Validator dependency	74
Simple validations on the bean	74
Custom validations	75
Unit testing	76
Controller setup	76
The Test method	76
<b>An overview of Spring MVC</b>	<b>77</b>
Important features	77
How it works	78
<b>Important concepts behind Spring MVC</b>	<b>79</b>
RequestMapping	79
Examples of request mapping	79
Example 1	79
Example 2	79
Example 3	80
Request Mapping methods - supported method arguments	81
RequestMapping methods - supported return types	82
View resolution	83
Configuring JSP view resolver	83

————— [ iii ] —————

Copyrighted material

Configuring Freemarker	84
Handler mappings and Interceptors	84
Defining a HandlerInterceptor	85
Mapping HandlerInterceptor to handlers	86
Model attributes	87
Session attributes	88

Putting an attribute in the session	88
Reading an attribute from the session	88
Removing an attribute from the session	89
<b>InitBinders</b>	89
The @ControllerAdvice annotation	89
<b>Spring MVC - advanced features</b>	90
Exception handling	90
Common exception handling across controllers	91
The error view	91
Specific exception handling in a Controller	92
Internationalization	93
Message bundle setup	94
Configuring a SessionLocaleResolver	94
Configuring a CookieLocaleResolver	95
Integration testing Spring controllers	96
Serving static resources	97
Exposing static content	97
Caching static content	98
Enabling GZip compression of static content	99
Integrating Spring MVC with Bootstrap	100
Bootstrap WebJar as Maven dependency	100
Configure Spring MVC resource handler to deliver WebJar static content	100
Using Bootstrap resources in JSP	101
<b>Spring Security</b>	101
Adding Spring Security dependency	102
Configuring a filter to intercept all requests	102
Logout	104
<b>Summary</b>	105
<b>Chapter 4: Evolution toward Microservices and Cloud-Native Applications</b>	106
<b>Typical web application architecture with Spring</b>	107
Web layer	108
Web application - rendering an HTML View	108
RESTful services	108
Business layer	108
Data layer	109

[ iv ]

Copyrighted material

Integration layer	109
Cross-cutting concerns	110
<b>Problems solved by Spring</b>	110
Loose coupling and testability	111
Plumbing code	111
Lightweight architecture	112
Architecture flexibility	112
Simplified implementation of cross-cutting concerns	112
Design patterns for free	112
<b>Application development goals</b>	113
Speed	114
Safety	114
Reliability	115
Availability	115
Security	115
Performance	116
High resilience	116
Scalability	116
<b>Challenges with monolithic applications</b>	117
Long release cycles	117
Difficult to scale	118
Adapting new technologies	118
Adapting new methodologies	118
Adapting modern development practices	118
<b>Understanding microservices</b>	118
What is a microservice?	119
The microservice architecture	119
Microservice characteristics	121
Small and lightweight microservices	121
Interoperability with message-based communication	122
Capability-aligned microservices	122
Independently deployable units	122
Stateless	122
Automated build and release process	123
Event-driven architecture	123
Approach 1 - sequential approach	124
Approach 2 - event-driven approach	124
Independent teams	125
Microservice advantages	125
Faster time to market	125
Technology evolution	126
Availability and scaling	126
Team dynamics	126
	127

[ v ]

Copyrighted material

<b>Microservice challenges</b>	127
Increased need for automation	127
Defining the boundaries of subsystems	128
Visibility and monitoring	128
Fault tolerance	129
Eventual consistency	129
Shared capabilities (enterprise level)	129
Increased need for operations teams	130
<b>Cloud-Native applications</b>	130
Twelve-Factor App	131
Maintain one code base	131
Dependencies	131
Config	133
Backing services	134
Build, release, run	134
Stateless	135
Port binding	135
Concurrency	135
Disposable	135
Environment parity	136
Logs as event streams	136
No distinction of admin processes	136
<b>Spring projects</b>	136
Spring Boot	137
Spring Cloud	137
<b>Summary</b>	138
<b>Chapter 5: Building Microservices with Spring Boot</b>	139
<b>What is Spring Boot?</b>	140
Building a quick prototype for a microservice	140
Primary goals	141
Nonfunctional features	141
<b>Spring Boot Hello World</b>	142

Configure spring-boot-starter-parent	142
spring-boot-starter-parent	144
Configure pom.xml with the required starter projects	145
Understanding starter projects	146
Configuring spring-boot-maven-plugin	148
Creating your first Spring Boot launch class	148
SpringApplication class	149
The @SpringBootApplication annotation	149
Running our Hello World application	150
Auto-configuration	152
Starter projects	156

— [ vi ] —

Copyrighted material

<b>What is REST?</b>	158
<b>First REST service</b>	160
Simple method returning string	161
Unit testing	161
Integration testing	162
Simple REST method returning an object	164
Executing a request	164
Unit testing	165
Integration testing	165
Get method with path variables	165
Executing a request	166
Unit testing	167
Integration testing	167
<b>Creating a todo resource</b>	168
Request methods, operations, and URIs	168
Beans and services	169
Retrieving a Todo list	171
Executing the service	171
Unit testing	172
Integration testing	173
Retrieving details for a specific Todo	174
Executing the service	175
Unit testing	175
Integration testing	176
Adding a Todo	176
Postman	177
Executing the POST service	178
Unit testing	179
Integration testing	180
<b>Spring Initializr</b>	181
Creating your first Spring Initializr project	183
pom.xml	185
FirstSpringInitializrApplication.java class	186
FirstSpringInitializrApplicationTests class	186
<b>A quick peek into auto-configuration</b>	186
<b>Summary</b>	189
<b>Chapter 6: Extending Microservices</b>	190
<b>Exception handling</b>	190
Spring Boot default exception handling	190
Nonexistent resource	191
Resource throwing an exception	192
Throwing a custom exception	193
Customizing the exception message	194

— [ vii ] —

Copyrighted material

Response status	196
<b>HATEOAS</b>	197
Sending HATEOAS links in response	199
Spring Boot starter HATEOAS	199
<b>Validation</b>	200
Enabling validation on the controller method	201
Defining validations on the bean	202
Unit testing validations	203
<b>Documenting REST services</b>	203
Generating a Swagger specification	204
Swagger UI	208
Customizing Swagger documentation using annotations	210
<b>Securing REST services with Spring Security</b>	213
Adding Spring Security starter	214
Basic authentication	214
Integration testing	216
Unit testing	217
OAuth 2 authentication	217
High-level flow	218
Implementing OAuth 2 authentication for our service	218
Setting up authorization and resource servers	218
Executing OAuth requests	219
Obtaining an access token	220
Executing the request using the access token	221
Integration test	222
Internationalization	223
<b>Caching</b>	226
Spring-boot-starter-cache	226
Enabling caching	226
Caching data	227
JSR-107 caching annotations	227
Auto-detection order	228
<b>Summary</b>	228
<b>Chapter 7: Advanced Spring Boot Features</b>	229
<b>Externalised configuration</b>	230
Customizing frameworks through application.properties	231
Logging	231
Embedded server configuration	232
Spring MVC	233
Spring Starter security	233
Data Sources, JDBC and JPA	233
Other configuration options	234

— [ viii ] —

Copyrighted material

Custom properties in application.properties	235
Configuration properties - type-safe Configuration Management	236

Profiles	239
Profiles-based Bean configuration	240
Other options for application configuration values	240
YAML configuration	241
Embedded servers	243
Switching to Jetty and Undertow	245
Building a WAR file	247
Developer tools	247
Live reload	248
Spring Boot Actuator	249
HAL Browser	251
Configuration properties	252
Environment details	254
Health	255
Mappings	255
Beans	257
Metrics	258
Auto-configuration	259
Debugging	261
Deploying an application to Cloud	261
Cloud Foundry	261
Summary	264
<b>Chapter 8: Spring Data</b>	265
Background - data stores	265
Spring Data	266
<b>Spring Data Commons</b>	266
Repository	267
The CrudRepository interface	267
The PagingAndSortingRepository interface	268
<b>Spring Data JPA</b>	268
Spring Data JPA example	269
New project with Starter Data JPA	269
Entities	270
The SpringApplication class	272
Populating some data	274
A simple repository	274
Unit test	275
The CrudRepository interface	276
Unit test	276

[ ix ]

Copyrighted material

The PagingAndSortingRepository interface	278
Unit tests	278
Query methods	280
Queries	282
Named Parameters	283
Named Query	283
Native query	283
<b>Spring Data Rest</b>	284
The GET method	286
The POST method	286
The search resource	288
<b>Big Data</b>	288
MongoDB	289
Unit test	290
Summary	291
<b>Chapter 9: Spring Cloud</b>	292
Introducing Spring Cloud	292
Spring Cloud Netflix	294
Demographic microservices setup	295
Microservice A	296
Service consumer	299
Ports	300
Centralized microservice configuration	301
Problem statement	301
Solution	302
Options	303
Spring Cloud Config	303
Implementing Spring Cloud Config Server	304
Setting up Spring Cloud Config Server	305
Connecting Spring Cloud Config Server to a local Git repository	306
Creating an environment-specific configuration	307
Spring Cloud Config Client	309
<b>Spring Cloud Bus</b>	311
The need for Spring Cloud Bus	311
Propagating configuration changes using Spring Cloud Bus	311
Implementation	312
<b>Declarative REST Client - Feign</b>	314
Load balancing	317
Ribbon	317
Implementation	318
The Name server	320

[ x ]

Copyrighted material

Limitations of hard coding microservice URLs	321
Workings of Name server	321
Options	322
Implementation	322
Setting up a Eureka Server	322
Registering microservices with Eureka	324
Connecting the service consumer microservice with Eureka	326
API Gateways	327
Implementing client-side load balancing with Zuul	328
Setting up a new Zuul API Gateway Server	329
Zuul custom filters	330
Invoking microservices through Zuul	332
Configuring service consumer to use Zuul API gateway	333
Distributed tracing	334
Distributed tracing options	335
Implementing Spring Cloud Sleuth and Zipkin	335
Integrating microservice components with Spring Cloud Sleuth	336
Setting up Zipkin Distributed Tracing Server	338
Integrating microservice components with Zipkin	339
Hystrix - fault tolerance	342
Implementation	342
Summary	343
<b>Chapter 10: Spring Cloud Data Flow</b>	344
Message-based asynchronous communication	344
Complexities of asynchronous communication	346
Spring projects for asynchronous messages	348
Sonata Integration	348

Spring Cloud Stream	349
Spring Cloud Data Flow	350
<b>Spring Cloud Stream</b>	351
Spring Cloud Stream architecture	352
Event processing - stock trading example	353
Model for stock trading example	354
The source application	354
Processor	356
Sink	358
<b>Spring Cloud Data Flow</b>	359
High-level architecture	361
Implementing Spring Cloud Data Flow	361
Setting up Spring Cloud Data Flow server	361
Setting up Data Flow Shell project	364
Configuring the apps	367

[ xi ]

Copyrighted material

Configuring the stream	369
Deploying the stream	370
Log messages - setting up connections to the message factory	371
Log messages - the flow of events	373
<b>Spring Cloud Data Flow REST APIs</b>	373
<b>Spring Cloud Task</b>	375
<b>Summary</b>	377
<b>Chapter 11: Reactive Programming</b>	378
<b>The Reactive Manifesto</b>	378
Characteristics of Reactive Systems	379
<b>Reactive use case - a stock price page</b>	380
The traditional approach	381
The reactive approach	381
Comparison between the traditional and reactive approaches	382
<b>Reactive programming in Java</b>	384
Reactive streams	384
Reactor	385
Mono	386
Flux	389
<b>Spring Web Reactive</b>	390
Creating a project using Spring Initializr	391
Creating an HTML view	394
Launching SpringReactiveExampleApplication	395
<b>Reactive databases</b>	397
Integrating MongoDB	398
Root Reactive MongoDB Starter	398
Creating a model object - a stock document	399
Creating a ReactiveCrudRepository	399
Initialising stock data using the Command Line Runner	400
Creating Reactive methods in Rest Controller	400
Updating the view to subscribe to the event stream	401
Launching SpringReactiveExampleApplication	402
<b>Summary</b>	402
<b>Chapter 12: Spring Best Practices</b>	404
<b>Maven standard directory layout</b>	404
<b>Layered architecture</b>	406
Recommended practices	407
Separate API and impl for important layers	408
<b>Exception handling</b>	409
Spring's approach to exception handling	410
The recommended approach	410
<b>Keeping your Spring configuration light</b>	411

[ xii ]

Copyrighted material

Using the basePackageClasses attribute in ComponentScan	411
Not using version numbers in schema references	411
Prefering constructor injection over setter injection for mandatory dependencies	412
<b>Managing dependency versions for Spring Projects</b>	413
<b>Unit testing</b>	415
The business layer	415
Web layer	415
The data layer	416
Other best practices	417
<b>Integration testing</b>	417
Spring Session	419
Example	420
Adding dependencies for Spring Session	420
Configuring Filter to replacing HttpSession with Spring Session	421
Enabling filtering for Tomcat by extending AbstractHttpSessionApplicationInitializer	421
<b>Caching</b>	422
Adding the Spring Boot Starter Cache dependency	422
Adding caching annotations	422
<b>Logging</b>	423
Logback	423
Log4j2	424
Framework independent configuration	425
<b>Summary</b>	425
<b>Chapter 13: Working with Kotlin in Spring</b>	426
<b>Kotlin</b>	426
<b>Kotlin versus Java</b>	427
Variables and type inference	428
Variables and immutability	428
Type system	429
Functions	430
Arrays	431
Collections	432
No checked exceptions	433
Data class	433
<b>Creating a Kotlin project in Eclipse</b>	434
Kotlin plugin	435
Creating a Kotlin project	437
Creating a Kotlin class	438

[ xiii ]

Copyrighted material

Running a Kotlin class	441
<b>Creating a Spring Boot project using Kotlin</b>	442
Dependencies and plugins	443
Spring Boot application class	445
Spring Boot application test class	446
<b>Implementing a REST service using Kotlin</b>	446
Simple method returning a string	447
Unit testing	447
Integration testing	448
Simple REST method returning an object	449
Executing a request	450
Unit testing	451
Integration testing	451
Get method with path variables	452
Executing a request	452
Unit testing	453
Integration testing	453
<b>Summary</b>	454
<b>Index</b>	455

## Preface

Spring 5.0 is due to arrive with a myriad of new and exciting features that will change the way we've used the framework so far. This book will show you this evolution—from solving the problems of testable applications to building distributed applications on the Cloud. The book begins with an insight into the new features in Spring 5.0, and shows you how to build an application using Spring MVC. You will then get a thorough understanding of how to build and extend microservices using the Spring Framework. You will also understand how to build and deploy Cloud applications. You will realize how application architectures have evolved from monoliths to those built around microservices. The advanced features of Spring Boot will also be covered and displayed through powerful examples.

By the end of this book, you will be equipped with the knowledge and best practices to develop applications with the Spring Framework.

### What this book covers

Chapter 1, *Evolution to Spring Framework 5.0*, takes you through the evolution of the Spring Framework, ranging from its initial versions to Spring 5.0. Initially, Spring was used to develop testable applications using dependency injection and core modules. Recent Spring Projects, such as Spring Boot, Spring Cloud, Spring Cloud Data Flow—deal with application infrastructure and moving applications to Cloud. We get an overview of different Spring modules and projects.

Chapter 2, *Dependency Injection*, dives deep into dependency injection. We will look at different kinds of dependency injection methods available in Spring, and how auto-wiring makes your life easy. We will also take a quick look into unit testing.

Chapter 3, *Building a Web Application with Spring MVC*, gives a quick overview of building a web application with Spring MVC.

Chapter 4, *Evolution toward Microservices and Cloud-Native Applications*, explains the evolution of application architectures in the last decade. We will understand why microservices and Cloud Native applications are needed and get a quick overview of the different Spring projects that help us build Cloud-Native applications.

---

### Preface

Chapter 5, *Building Microservices with Spring Boot*, discusses how Spring Boot takes away the complexity in creating production-grade Spring-based applications. It makes it easy to get started with Spring-based projects and provides easy integration with third-party libraries. In this chapter, we will take the students on a journey with Spring Boot. We will start with implementing a basic web service and then move on to adding caching, exception handling, HATEOAS, and internationalization, while making use of different features from the Spring Framework.

Chapter 6, *Extending Microservices*, focuses on adding more advanced features to the microservices that we built in Chapter 4, *Evolution toward Microservices and Cloud-Native Applications*.

Chapter 7, *Advanced Spring Boot Features*, takes a look at the advanced features in Spring Boot. You will learn how to monitor a microservice with a Spring Boot Actuator. Then, you will deploy the microservice to Cloud. You will also learn how to develop more effectively with the developer tools provided by Spring Boot.

Chapter 8, *Spring Data*, discusses the Spring Data module. We will develop simple applications to integrate Spring with JPA and Big Data technologies.

Chapter 9, *Spring Cloud*, discusses the distributed systems in the Cloud that have common problems, configuration management, service discovery, circuit breakers, and intelligent routing. In this chapter, you will learn how Spring Cloud helps you develop solutions for these common problems. These solutions should work well on the Cloud as well as

Chapter 10, *Spring Cloud Data Flow*, talks about the Spring Cloud Data Flow, which offers a collection of patterns and best practices for microservices-based distributed streaming and batch data pipelines. In this chapter, we will understand the basics of Spring Cloud Data Flow and use it to build basic data flow use cases.

Chapter 11, *Reactive Programming*, explores programming with asynchronous data streams. In this chapter, we will understand Reactive programming and take a quick look at the features provided by the Spring framework.

Chapter 12, *Spring Best Practices*, helps you understand best practices in developing enterprise applications with Spring related to unit testing, integration testing, maintaining Spring configuration, and more.

Chapter 13, *Working with Kotlin in Spring*, introduces you to a JVM language gaining quick popularity—Kotlin. We will discuss how to setup a Kotlin project in Eclipse. We will create a new Spring Boot project using Kotlin and implement a couple of basic services with unit and integration testing.

## What you need for this book

To be able to run examples from this book, you will need the following tools:

- Java 8
- Eclipse IDE
- Postman

We will use Maven embedded into Eclipse IDE to download all the dependencies that are needed.

## Who this book is for

This book is for experienced Java developers who knows the basics of Spring, and wants to learn how to use Spring Boot to build applications and deploy them to the Cloud.

## Conventions

In this book, you will find a number of text styles that distinguish between different kinds of information. Here are some examples of these styles and an explanation of their meaning.

Code words in text, database table names, folder names, filenames, file extensions, pathnames, dummy URLs, and user input are shown as follows: "Configure `spring-boot-starter-parent` in your `pom.xml` file".

A block of code is set as follows:

```
<properties>
  <mockito.version>1.10.20</mockito.version>
</properties>
```

Any command-line input or output is written as follows:

```
mvn clean install
```

New terms and important words are shown in bold. Words that you see on the screen, for example, in menus or dialog boxes, appear in the text like this: "Provide the details and click on **Generate Project**."



Warnings or important notes appear in a box like this.



Tips and tricks appear like this.



## Reader feedback

Feedback from our readers is always welcome. Let us know what you think about this book—what you liked or disliked. Reader feedback is important for us as it helps us develop titles that you will really get the most out of.

To send us general feedback, simply e-mail [feedback@packtpub.com](mailto:feedback@packtpub.com), and mention the book's title in the subject of your message.

If there is a topic that you have expertise in and you are interested in either writing or contributing to a book, see our author guide at [www.packtpub.com/authors](http://www.packtpub.com/authors).

## Customer support

Now that you are the proud owner of a Packt book, we have a number of things to help you to get the most from your purchase.

## Downloading the example code

You can download the example code files for this book from your account at <http://www.packtpub.com>. If you purchased this book elsewhere, you can visit <http://www.packtpub.com/support> and register to have the files e-mailed directly to you.

You can download the code files by following these steps:

1. Log in or register to our website using your e-mail address and password.
2. Hover the mouse pointer on the SUPPORT tab at the top.
3. Click on **Code Downloads & Errata**.
4. Enter the name of the book in the **Search** box.
5. Select the book for which you're looking to download the code files.

6. Choose from the drop-down menu where you purchased this book from.
7. Click on **Code Download**.

Once the file is downloaded, please make sure that you unzip or extract the folder using the latest version of:

- WinRAR / 7-Zip for Windows
- Zipper / iZip / UnRarX for Mac
- 7-Zip / PeaZip for Linux

The code bundle for the book is also hosted on GitHub at <https://github.com/PacktPublishing/Mastering-Spring-5.0>. We also have other code bundles from our rich catalog of books and videos available at <https://github.com/PacktPublishing/>. Check them out!

## Errata

Although we have taken every care to ensure the accuracy of our content, mistakes do happen. If you find a mistake in one of our books—maybe a mistake in the text or the code—we would be grateful if you could report this to us. By doing so, you can save other readers from frustration and help us improve subsequent versions of this book. If you find any errata, please report them by visiting <http://www.packtpub.com/submit-errata>, selecting your book, clicking on the **Errata Submission Form** link, and entering the details of your errata. Once your errata are verified, your submission will be accepted and the errata will be uploaded to our website or added to any list of existing errata under the Errata section of that title.

To view the previously submitted errata, go to <https://www.packtpub.com/books/content/support> and enter the name of the book in the search field. The required information will appear under the **Errata** section.

## Piracy

Piracy of copyrighted material on the Internet is an ongoing problem across all media. At Packt, we take the protection of our copyright and licenses very seriously. If you come across any illegal copies of our works in any form on the Internet, please provide us with the location address or website name immediately so that we can pursue a remedy.

Please contact us at [copyright@packtpub.com](mailto:copyright@packtpub.com) with a link to the suspected pirated material. We appreciate your help in protecting our authors and our ability to bring you valuable content.

## Questions

If you have a problem with any aspect of this book, you can contact us at [questions@packtpub.com](mailto:questions@packtpub.com), and we will do our best to address the problem.

We will want to create a mock for `DataService`. There are multiple approaches to creating mocks with Mockito. Let's use the simplest among them—annotations. We use the `@Mock` annotation to create a mock for `DataService`:

```
@Mock  
private DataService DataService;
```

Once we create the mock, we will need to inject it into the class under test, `BusinessServiceImpl`. We do that using the `@InjectMocks` annotation:

```
@InjectMocks  
private BusinessService service =  
new BusinessServiceImpl();
```

In the test method, we will need to stub the mock service to provide the data that we want it to provide. There are multiple approaches. We will use the `RDD` stub methods provided by

Copyrighted material

Mockito to mock the `retrieveData` method:

```
BDDMockito.given(dataService.retrieveData(  
    Matchers.any(User.class))  
    .willReturn(Arrays.asList(new Data(10),  
    new Data(15), new Data(25))));
```

What we are defining in the preceding code is called **stubbing**. As with anything with Mockito, this is extremely readable. When the `retrieveData` method is called on the `dataService` mock with any object of type `User`, it returns a list of three items with values specified.

When we use Mockito annotations, we would need to use a specific JUnit runner, that is, `MockitoJUnitRunner`. `MockitoJUnitRunner` helps in keeping the test code clean and provides clear debugging information in case of test failures. `MockitoJUnitRunner` initializes the beans annotated with `@Mock` annotation and also validates the usage of framework after execution of each test method.

```
@RunWith(MockitoJUnitRunner.class)
```

The complete list of the test is as follows. It has one test method:

```
@RunWith(MockitoJUnitRunner.class)  
public class BusinessServiceMockitoTest {  
    private static final User DUMMY_USER = new User("dummy");  
    @Mock  
    private DataService dataService;  
    @InjectMocks  
    private BusinessService service =  
        new BusinessServiceImpl();
```

[ 41 ]

Copyrighted material

#### Dependency Injection

```
@Test  
public void testCalculateSum() {  
    BDDMockito.given(dataService.retrieveData(  
        Matchers.any(User.class)))  
        .willReturn(  
            Arrays.asList(new Data(10),  
            new Data(15), new Data(25)));  
    long sum = service.calculateSum(DUMMY_USER);  
    assertEquals(10 + 15 + 25, sum);  
}
```

### Container managed beans

Instead of a class creating its own dependencies, in the earlier example, we looked at how the Spring IoC container can take over the responsibility of managing beans and their dependencies. The beans that are managed by the container are called **Container Managed Beans**.



[ 42 ]

Copyrighted material

#### Dependency Injection

Delegating the creation and management of beans to the container has many advantages. Some of them are listed as follows:

- Since classes are not responsible for creating dependencies, they are loosely coupled and testable. This leads to good design and fewer defects.
- Since the container manages the beans, a few hooks around the beans can be introduced in a more generic way. Cross-cutting concerns, such as logging, caching, transaction management, and exception handling can be woven around these beans using **Aspect-Oriented Programming (AOP)**. This leads to more maintainable code.

### Dependency injection types

In the previous example, we used a setter method to wire in the dependency. There are two types of dependency injections that are used frequently:

- The setter injection
- The constructor injection

#### The setter injection

The setter injection is used to inject the dependencies through setter methods. In the following example, the instance of `DataService` uses the setter injection:

```
public class BusinessServiceImpl {  
    private DataService dataService;  
    @Autowired  
    public void setDataService(DataService dataService) {  
        this.dataService = dataService;  
    }  
}
```

Actually, in order to use the setter injection, you do not even need to declare a setter method. If you specify `@Autowired` on the variable, Spring automatically uses the setter injection. So, the following code is all that you need for the setter injection for `DataService`:

```
public class BusinessServiceImpl {  
    @Autowired
```

```
    private DataService dataService;  
}
```

[ 43 ]

Copyrighted material

#### Dependency Injection

### The constructor injection

The constructor injection, on the other hand, uses a constructor to inject dependencies. The following code shows how to use a constructor for injecting in `DataService`:

```
public class BusinessServiceImpl {  
    private DataService dataService;  
    @Autowired  
    public BusinessServiceImpl(DataService dataService) {  
        super();  
        this.dataService = dataService;  
    }  
}
```

When you run the code with the preceding implementation of `BusinessServiceImpl`, you will see this statement in the log, asserting that autowiring took place using the constructor:

```
Autowiring by type from bean name 'businessServiceImpl' via  
constructor to bean named 'dataServiceImpl'
```

### Constructor versus setter injection

Originally, in XML-based application contexts, we used the constructor injection with mandatory dependencies and the setter injection with nonmandatory dependencies.

However, an important thing to note is that when we use `@Autowired` on a field or a method, the dependency is required by default. If no candidates are available for an `@Autowired` field, autowiring fails and throws an exception. So, the choice is not so clear anymore with Java application contexts.

Using the setter injection results in the state of the object changing during the creation. For fans of immutable objects, the constructor injection might be the way to go. Using the setter injection might sometimes hide the fact that a class has a lot of dependencies. Using the constructor injection makes it obvious, since the size of the constructor increases.

### Spring bean scopes

Spring beans can be created with multiple scopes. The default scope is a singleton.

Since there is only one instance of a singleton bean, it cannot contain any data that is specific to a request.

[ 44 ]

Copyrighted material

Pages 45 to 65 are not shown in this preview.

#### Building a Web Application with Spring MVC

One thing to note is this:

- `$(name)`: This uses the **Expression Language (EL)** syntax to access the attribute from the model.

Here is a screenshot of how this would look on the screen when the URL is hit:



### Unit testing

In this unit test, we want to set up `BasicModelMapController`, execute a get request to `/welcome-model-map`, and check whether the model has the expected attribute and whether the expected view name is returned.

#### Setting up the Controller to test

This step is very similar to the previous flow. We instantiate Mock MVC with `BasicModelMapController`:

```
this.mockMvc = MockMvcBuilders.standaloneSetup  
(new BasicModelMapController())  
.setViewResolvers(viewResolver()).build();
```

#### Writing the Test method

The complete Test method is shown in the following code:

```
@Test  
public void basicTest() throws Exception {  
    this.mockMvc  
.perform()
```

[ 66 ]

Copyrighted material

#### Building a Web Application with Spring MVC

```
get("/welcome-model-map")  
.accept(MediaType.parseMediaType  
("application/html;charset=UTF-8"))  
.andExpect(model().attribute("name", "XYZ"))  
.andExpect(view().name("welcome-model-map"));  
}
```

A few important things to note:

- `get("/welcome-model-map")`: Execute `get` request to the specified URL
- `model().attribute("name", "XYZ")`: Result Matcher to check if the model contains specified attribute `name` with specified value `XYZ`
- `view().name("welcome-model-map")`: Result Matcher to check if the view name returned is as specified

## Flow 4 - Controller redirecting to a View with ModelAndView

In the previous flow, we returned a view name and populated the model with attributes to be used in the view. Spring MVC provides an alternate approach using `ModelAndView`. The controller method can return a `ModelAndView` object with the view name and appropriate attributes in the Model. In this flow, we will explore this alternate approach.

### Spring MVC controller

Take a look at the following controller:

```
@Controller
public class BasicModelViewController {
    @RequestMapping(value = "/welcome-model-view")
    public ModelAndView welcome(ModelMap model) {
        model.put("name", "XYZ");
        return new ModelAndView("welcome-model-view", model);
    }
}
```

[ 67 ]

Copyrighted material

*Building a Web Application with Spring MVC*

A few important things to note are as follows:

- `@RequestMapping(value = "/welcome-model-view")`: The URI mapped is `/welcome-model-view`.
- `public ModelAndView welcome(ModelMap model)`: Note that the return value is no longer a String. It is `ModelAndView`.
- `return new ModelAndView("welcome-model-view", model)`: Create a `ModelAndView` object with the appropriate view name and model.

### Creating a View

Let's create a view using the model attribute `name` that was set in the model in the controller. Create a simple JSP in the `/WEB-INF/views/welcome-model-view.jsp` path:

Welcome \${name}! This is coming from a model-view - a JSP

Here is a screenshot of how this would look on the screen when the URL is hit:



### Unit testing

Unit testing for this flow is similar to the previous flow. We would need to check if the expected view name is returned.

[ 68 ]

Copyrighted material

*Building a Web Application with Spring MVC*

## Flow 5 - Controller redirecting to a View with a form

Now let's shift our attention to creating a simple form to capture input from the user.

The following steps will be needed:

- Create a simple POJO. We want to create a user. We will create a POJO User.
- Create a couple of Controller methods—one to display the form, and the other to capture the details entered in the form.
- Create a simple View with the form.

### Creating a command or form backing object

POJO stands for Plain Old Java Object. It is usually used to represent a bean following the typical JavaBean conventions. Typically, it contains private member variables with getters and setters and a no-argument constructor.

We will create a simple POJO to act as a command object. Important parts of the class are listed as follows:

```
public class User {
    private String guid;
    private String name;
    private String userId;
    private String password;
    private String password2;
    //Constructor
    //Getters and Setters
    //toString
}
```

A few important things to note are as follows:

- This class does not have any annotations or Spring-related mappings. Any bean can act as a form-backing object.
- We are going to capture the `name`, `user ID`, and `password` in the form. We have a password confirmation field, `password2`, and unique identifier field `guid`.
- Constructor, getters, setters, and `toString` methods are not shown for brevity.

#### *Building a Web Application with Spring MVC*

- **Implicit determination of the View:** If a view name is not present in the return type, it is determined using `DefaultRequestToViewNameTranslator`. By default, `DefaultRequestToViewNameTranslator` removes the leading and trailing slashes as well as the file extension from the URI; for example, the `display.html` becomes `display`.

The following are some of the return types that are supported on Controller methods with Request Mapping:

Return Type	What happens?
ModelAndView	The object includes a reference to the model and the view name.
Model	Only Model is returned. The view name is determined using <code>DefaultRequestToViewNameTranslator</code> .
Map	A simple map to expose a model.
View	A view with a model implicitly defined.
String	Reference to a view name.

## View resolution

Spring MVC provides very flexible view resolution. It provides multiple view options:

- Integration with JSP, Freemarker.
- Multiple view resolution strategies. A few of them are listed as follows:
  - `XmlViewResolver`: View resolution based on an external XML configuration
  - `ResourceBundleViewResolver`: View resolution based on a property file
  - `UrlBasedViewResolver`: Direct mapping of the logical view name to a URL
  - `ContentNegotiatingViewResolver`: Delegates to other view resolvers based on the Accept request header
- Support for chaining of view resolvers with the explicitly defined order of preference.
- Direct generation of XML, JSON, and Atom using Content Negotiation.

#### *Building a Web Application with Spring MVC*

### Configuring JSP view resolver

The following example shows the commonly used approach to configure a JSP view resolver using `InternalResourceViewResolver`. The physical view name is determined using the configured prefix and suffix for the logical view name using `JstlView`:

```
<bean id="jspViewResolver" class="org.springframework.web.servlet.view.InternalResourceViewResolver">
<property name="viewClass"
  value="org.springframework.web.servlet.view.JstlView"/>
<property name="prefix" value="/WEB-INF/jsp/"/>
<property name="suffix" value=".jsp"/>
</bean>
```

There are other approaches using property and XML files for mapping.

### Configuring Freemarker

The following example shows the typical approach used to configure a Freemarker view resolver.

First, the `freemarkerConfig` bean is used to load the Freemarker templates:

```
<bean id="freemarkerConfig"
  class="org.springframework.web.servlet.view.freemarker.FreeMarkerConfigurer">
<property name="templateLoaderPath" value="/WEB-INF/freemarker"/>
</bean>
```

The following bean definition shows how to configure a Freemarker view resolver:

```
<bean id="freemarkerViewResolver"
  class="org.springframework.web.servlet.view.freemarker.FreeMarkerViewResolver">
<property name="cache" value="true"/>
<property name="prefix" value="/" />
<property name="suffix" value=".ftl"/>
</bean>
```

As with JSPs, the view resolution can be defined using properties or an XML file.

## Handler mappings and Interceptors

In the version before Spring 2.5 (before there was support for Annotations), the mapping between a URL and a Controller (also called a handler) was expressed using something called a handler mapping. It is almost a historical fact today. The use of annotations eliminated the need for an explicit handler mapping.

HandlerInterceptors can be used to intercept requests to handlers (or controllers). Sometimes, you would want to do some processing before and after a request. You might want to log the content of the request and response, or you might want to find out how much time a specific request took.

There are two steps in creating a HandlerInterceptor:

1. Define the HandlerInterceptor.
2. Map the HandlerInterceptor to the specific handlers to be intercepted.

### Defining a HandlerInterceptor

The following are the methods you can override in HandlerInterceptorAdapter:

- public boolean preHandle(HttpServletRequest request, HttpServletResponse response, Object handler): Invoked before the handler method is invoked
- public void postHandle(HttpServletRequest request, HttpServletResponse response, Object handler, ModelAndView modelAndView): Invoked after the handler method is invoked
- public void afterCompletion(HttpServletRequest request, HttpServletResponse response, Object handler, Exception ex): Invoked after the request processing is complete

The following example implementation shows how to create a HandlerInterceptor. Let's start with creating a new class that extends HandlerInterceptorAdapter:

```
public class HandlerTimeLoggingInterceptor extends
HandlerInterceptorAdapter {
```

[ 85 ]

Copyrighted material

Pages 86 to 96 are not shown in this preview.

A few things to note are as follows:

- @RunWith(SpringRunner.class): SpringRunner helps us launch a Spring context.
- @WebAppConfiguration: Used to launch a web app context with Spring MVC
- @ContextConfiguration("file:src/main/webapp/WEB-INF/user-web-context.xml"): Specifies the location of the spring context XML.
- this.mockMvc = MockMvcBuilders.webAppContextSetup(this.wac).build(): In the earlier examples, we used a standalone setup. However, in this example, we want to launch the entire web app. So, we use webAppContextSetup.
- The execution of the test is very similar to how we did it in earlier tests.

## Serving static resources

Most teams today have separate teams delivering frontend and backend content. The frontend is developed with modern JavaScript frameworks, such as AngularJS, Backbone, and so on. Backend is built through web applications or REST services based on frameworks such as Spring MVC.

With this evolution in frontend frameworks, it is very important to find the right solutions to version and deliver frontend static content.

The following are some of the important features provided by the Spring MVC framework:

- They expose static content from folders in the web application root
- They enable caching
- They enable GZip compression of static content

## Exposing static content

Web applications typically have a lot of static content. Spring MVC provides options to expose static content from folders on the web application root as well locations on the classpath. The following snippet shows that content within the war can be exposed as static content:

```
<mvc:resources
mapping="/resources/**"
location="/static-resources"/>
```

[ 97 ]

Copyrighted material

Things to note are as follows:

- location="/static-resources/": The location specifies the folders inside the war or classpath that you would want to expose as static content. In this example, we want to expose all the content in the static-resources folder inside the root of war as static content. We can specify multiple comma-separated values to expose multiple folders under the same external facing URL.
- mapping="/resources/\*\*": The mapping specifies the external facing URI path. So, a CSS file named app.css inside the static-resources folder can be accessed using the /resources/app.css URL.

The complete Java configuration for the same configuration is shown here:

```
@Configuration
@EnableWebMvc
public class WebConfig extends WebMvcConfigurerAdapter {
    @Override
    public void addResourceHandlers
```

```

(ResourceHandlerRegistry registry) {
    registry
        .addResourceHandler("/*")
        .addResourceLocations("/*");
}
}

```

### Caching static content

Caching for static resources can be enabled for improved performance. The browser would cache the resources served for the specified time period. The `cache-period` attribute or the `setCachePeriod` method can be used to specify the caching interval (in seconds) based on the type of configuration used. The following snippets show the details:

This is the Java configuration:

```

registry
    .addResourceHandler("/*")
    .addResourceLocations("/*");
    .setCachePeriod(365 * 24 * 60 * 60);
}
}

```

This is the XML configuration:

```

<mvc:resources
    mapping="/*"
    location="static-resources/"
    cache-period="365 * 24 * 60 * 60"/>

```

[ 98 ]

Copyrighted material

### *Building a Web Application with Spring MVC*

The `Cache-Control: max-age={specified-max-age}` response header will be sent to the browser.

### Enabling GZip compression of static content

Compressing a response is a simple way to make web applications faster. All modern browsers support GZip compression. Instead of sending the full static content file, a compressed file can be sent as a response. The browser will decompress and use the static content.

The browser can specify that it can accept the compressed content with a request header. If the server supports it, it can deliver the compressed content—again, marked with a response header.

Request Header sent from browser is as follows:

```
Accept-Encoding: gzip, deflate
```

Response Header sent from the web application is as follows:

```
Content-Encoding: gzip
```

The following snippet shows how to add a Gzip resolver to deliver compressed static content:

```

registry
    .addResourceHandler("/*")
    .addResourceLocations("/*");
    .setCachePeriod(365 * 24 * 60 * 60)
    .resourceChain(true)
    .addResolver(new GzipResourceResolver())
    .addResolver(new PathResourceResolver());
}
}

```

Things to note are as follows:

- `resourceChain(true)`: We would want to enable Gzip compression, but would want to fall back to delivering the full file if full file was requested. So, we use resource chaining (chaining of resource resolvers).
- `addResolver(new PathResourceResolver())`: `PathResourceResolver`: This is the default resolver. It resolves based on the resource handlers and locations configured.
- `addResolver(new GzipResourceResolver())`: `GzipResourceResolver`: This enables Gzip compression when requested.

[ 99 ]

Copyrighted material

Pages 100 to 102 are not shown in this preview.

### *Building a Web Application with Spring MVC*

```

<url-pattern>/*</url-pattern>
</filter-mapping>

```

Now, all requests to our web application will go through the filter. However, we have not configured anything related to security yet. Let's use a simple Java configuration example:

```

@Configuration
@EnableWebSecurity
public class SecurityConfiguration extends
WebSecurityConfigurerAdapter {
    @Autowired
    public void configureGlobalSecurity
    (AuthenticationManagerBuilder auth) throws Exception {
        auth
            .inMemoryAuthentication()
            .withUser("firstuser").password("password1")
            .roles("USER", "ADMIN");
    }
    @Override
    protected void configure(HttpSecurity http)
    throws Exception {
        http
            .authorizeRequests()
            .antMatchers("/login").permitAll()
            .antMatchers("/*secure**/*")
            .access("hasRole('USER')");
            .and().formLogin();
    }
}

```

Things to note are as follows:

- `@EnableWebSecurity`: This annotation enables any Configuration class to contain the definition of Spring Configuration. In this specific instance, we override a couple of methods to provide our specific Spring MVC configuration.
- `WebSecurityConfigurerAdapter`: This class provides a base class to create a Spring configuration (`WebSecurityConfigurer`).
- `protected void configure(HttpSecurity http)`: This method provides the security needs for different URLs.
- `antMatchers("/*secure**/*").access("hasRole('USER')")`: You would need a role of `USER` to access any URL containing the sub-string `secure`.

- `public void configureGlobalSecurity(AuthenticationManagerBuilder auth)`: In this example, we are using in-memory authentication. This can be used to connect to a database (`auth.jdbcAuthentication()`), or an LDAP (`auth.ldapAuthentication()`), or a custom authentication provider (created extending `AuthenticationProvider`).
- `.withUser("firstuser").password("password1")`: Configures an in-memory valid user ID and password combination.
- `.roles("USER", "ADMIN")`: Assigns roles to the user.

When we try to access any secure URLs, we will be redirected to a login page. Spring Security provides ways of customizing the Logic page as well as the redirection. Only authenticated users with the right roles will be allowed to access the secured application pages.

## Logout

Spring Security provides features to enable a user to log out and be redirected to a specified page. The URI of the `LogoutController` is typically mapped to the Logout link in the UI. The complete listing of `LogoutController` is as follows:

```
@Controller
public class LogoutController {
    @RequestMapping(value = "/secure/logout",
    method = RequestMethod.GET)
    public String logout(HttpServletRequest request,
    HttpServletResponse response) {
        Authentication auth =
        SecurityContextHolderHolder.getContext()
        .getAuthentication();
        if (auth != null) {
            new SecurityContextLogoutHandler()
            .logout(request, response, auth);
            request.getSession().invalidate();
        }
        return "redirect:/secure/welcome";
    }
}
```

Things to note are as follows:

- `if (auth != null):` If there is a valid authentication, then end the session
- `new SecurityContextLogoutHandler().logout(request, response, auth);`: `SecurityContextLogoutHandler` performs a logout by removing the authentication information from `SecurityContextHolder`
- `return "redirect:/secure/welcome"`: Redirects to the secure welcome page

## Summary

In this chapter, we discussed the basics of developing web applications with Spring MVC. We also discussed implementing exception handling, internationalization and securing our applications with Spring Security.

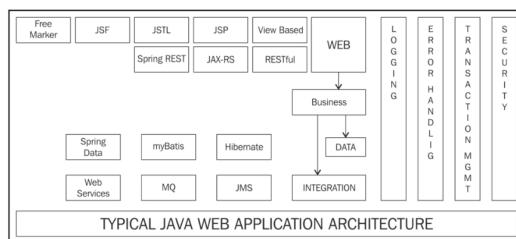
Spring MVC can also be used to build REST services. We will discuss that and more related to REST services in the subsequent chapters.

In the next chapter, we will shift our attention toward microservices. We will try to understand why the world is looking keenly at microservices. We will also explore the importance of applications being Cloud-Native.

- What are the challenges with microservices?
- What are the good practices that help in deploying microservices to the Cloud?
- What are the Spring projects that help us in developing microservices and Cloud-Native applications?

## Typical web application architecture with Spring

Spring has been the framework of choice to wire Java Enterprise applications during the last decade and half. Applications used a layered architecture with all cross-cutting concerns being managed using aspect-oriented programming. The following diagram shows a typical architecture for a web application developed with Spring:



The typical layers in such an application are listed here. We will list cross-cutting concerns as a separate layer, though in reality, they are applicable across all layers:

- **Web layer:** This is typically responsible for controlling the web application flow (controller and/or Front Controller) and rendering the view.
- **Business layer:** This is where all your business logic is written. Most applications have transaction management starting from the business layer.
- **Data layer:** It is also responsible for talking to the database. This is responsible for persisting/retrieving data in Java objects to the tables in the database.

[ 107 ]

Copyrighted material

### *Evolution toward Microservices and Cloud-Native Applications*

- **Integration layer:** Applications talk to other applications, either over queues or by invoking web services. The integration layer establishes such connections with other applications.
- **Cross-cutting concerns:** These are concerns across different layers—logging, security, transaction management, and so on. Since Spring IoC container manages the beans, it can weave these concerns around the beans through **Aspect-oriented programming (AOP)**.

Let's discuss each of the layers and the frameworks used in more detail.

## Web layer

Web layer is dependent on how you would want to expose the business logic to the end user. Is it a web application? Or are you exposing RESTful web services?

### Web application - rendering an HTML View

These web applications use a web MVC framework such as Spring MVC or Struts. The View can be rendered using JSP, JSF, or template-based frameworks such as Freemarker.

### RESTful services

There are two typical approaches used to develop RESTful web services:

- **JAX-RS:** The Java API for REST services. This is a standard from the Java EE specification. Jersey is the reference implementation.
- **Spring MVC or Spring REST:** Restful services can also be developed with Spring MVC.

Spring MVC does not implement JAX-RS so, the choice is tricky. JAX-RS is a Java EE standard. But Spring MVC is more innovative and more likely to help you build new features faster.

## Business layer

The business layer typically contains all the business logic in an application. Spring Framework is used in this layer to wire beans together.

[ 108 ]

Copyrighted material

### *Evolution toward Microservices and Cloud-Native Applications*

This is also the layer where the boundary of transaction management begins. Transaction management can be implemented using Spring AOP or AspectJ. A decade back, **Enterprise Java Beans (EJB)** were the most popular approach to implement your business layer. With its lightweight nature, Spring is now the framework of choice for the business layer.

EJB3 is much simpler than EJB2. However, EJB3 is finding it difficult to recover the ground lost to Spring.

## Data layer

Most applications talk to a database. The data layer is responsible for storing data from your Java objects to your database and vice versa. The following are the most popular approaches to building data layers:

- **JPA:** The Java Persistence API helps you to map Java objects (POJOs) to your database tables. Hibernate is the most popular implementation for JPA. JPA is typically preferred for all transactional applications. JPA is not the best choice for batch and reporting applications.
- **MyBatis:** MyBatis (previously, iBatis) is a simple data-mapping framework. As its website (<http://www.mybatis.org/mybatis-3/>) says, *MyBatis is a first class persistence framework with support for custom SQL, stored procedures and advanced mappings. MyBatis eliminates almost all of the JDBC code and manual setting of parameters and retrieval of results.* MyBatis can be considered for batch and reporting applications where SQLs and stored procedures are more typically used.
- **Spring JDBC:** JDBC and Spring JDBC are not that commonly used anymore.

We will discuss in detail the advantages and disadvantages of JDBC, Spring JDBC, MyBatis

## Integration layer

The integration layer is typically where we talk to other applications. There might be other applications exposing SOAP or RESTful services over HTTP (the web) or MQ.

- Spring JMS is typically used to send or receive messages on queues or service buses.
- Spring MVC RestTemplate can be used to invoke RESTful services.
- Spring WS can be used to invoke SOAP-based web services.

[ 109 ]

Copyrighted material

Page 110 is not part of this book preview.

## Problems solved by Spring

Spring is the framework of choice to wire Enterprise Java applications. It has solved a number of problems that Enterprise Java applications have faced since the complexity associated with EJB2. A few of them are listed as follows:

- Loose coupling and testability
- Plumbing code
- Lightweight architecture
- Architectural flexibility
- Simplified implementation of cross-cutting concerns
- Best design patterns for free

## Loose coupling and testability

Through dependency injection, Spring brings loose coupling between classes. While loose coupling is beneficial to application maintainability in the long run, the first benefits are realized with the testability that it brings in.

Testability was not a forte of Java EE (or J2EE, as it was called then) before Spring. The only way to test EJB2 applications was to run them in the container. Unit testing them was incredibly difficult.

That's exactly the problem Spring Framework set out to solve. As we saw in the earlier chapters, if objects are wired using Spring, writing unit tests becomes easier. We can easily stub or mock dependencies and wire them into objects.

## Plumbing code

The developer of the late 1990s and early-to-mid-2000s will be familiar with the amount of plumbing code that had to be written to execute a simple query through JDBC and populate the result into a Java object. You had to perform a [Java Naming and Directory Interface \(JNDI\)](#) lookup, get a connection, and populate the results. This resulted in duplicate code. Usually, problems were repeated with exception handling code in every method. And this problem is not limited to JDBC.

One of these problems Spring Framework solved was by eliminating all the plumbing code. With Spring JDBC, Spring JMS, and other abstractions, the developers could focus on writing business logic. Spring framework took care of the nitty-gritty.

[ 111 ]

Copyrighted material

## Lightweight architecture

Using EJBs made the applications complex, and not all applications needed that complexity. Spring provided a simplified, lightweight way of developing applications. If distribution was needed, it could be added later.

## Architecture flexibility

Spring Framework is used to wire objects across an application in different layers. In spite of its ever-loomng presence, Spring Framework did not restrict the flexibility or choice of frameworks that application architects and developers had. A couple of examples are listed as follows:

- Spring Framework provided great flexibility in the web layer. If you wanted to use Struts or Struts 2 instead of Spring MVC, it was configurable. You had the choice of integrating with a wider range of view and template frameworks.
- Another good example is the data layer, where you had possibilities to connect with JPA, JDBC, and mapping frameworks, such as MyBatis.

## Simplified implementation of cross-cutting concerns

When Spring Framework is used to manage beans, the Spring IoC container manages the life cycle—creation, use, auto-wiring, and destruction—of the beans. It makes it easier to weave an additional functionality—such as the cross-cutting concerns—around the beans.

## Design patterns for free

Spring Framework encourages the use of a number of design patterns by default. A few examples are as follows:

- **Dependency Injection or Inversion of Controller:** This is the fundamental design pattern Spring Framework is built to enable. It enables loose coupling and testability.
- **Singleton:** All Spring beans are singletons by default.
- **Factory Pattern:** Using the bean factory to instantiate beans is a good example of the factory pattern.

[ 112 ]

Copyrighted material

- **Front Controller:** Spring MVC uses DispatcherServlet as the Front Controller. So we use the Front Controller pattern when we develop applications with Spring MVC.
- **Template Method:** Helps us avoid boilerplate code. Many Spring-based classes—JdbcTemplate and JmsTemplate—are implementations of this pattern.

## Application development goals

Before we move on to the concepts of REST services, microservices, and Cloud-Native applications, let's take some time to understand the common goals we have when we develop applications. Understanding these goals will help us understand why applications are moving toward the microservices architecture.

First of all, we should remember that the software industry is still a relatively young industry. One thing that's been a constant in my decade and a half experience with developing, designing, and architecting software is that things change. The requirements of today are not the requirements of tomorrow. Technology today is not the technology we will use tomorrow. While we can try predicting what happens in the future, we are often wrong.

One of the things we did during the initial decades of software development was build software systems for the future. The design and architecture were made complex in preparation for future requirements.

During the last decade, with **agile** and **extreme programming**, the focus shifted to being **lean** and building good enough systems, adhering to basic principles of design. The focus shifted to evolutionary design. The thought process is this: *If a system has good design for today's needs, and is continuously evolving and has good tests, it can easily be refactored to meet tomorrow's needs.*

While we do not know where we are heading, we do know that a big chunk of our goals when developing applications have not changed.

The key goals of software development, for a large number of applications, can be described with the statement *speed and safety at scale*.

We will discuss each of these in elements in the next section.

[ 113 ]

Copyrighted material

## Speed

The speed of delivering new requirements and innovations is increasingly becoming a key differentiator. It is not sufficient to develop (code and test) fast. It is important to deliver (to production) quickly. It is now common knowledge that the best software organizations in the world deliver software to production multiple times every day.

The technology and business landscape is in a constant flux, and is constantly evolving. The key question is "How fast can an application adapt to these changes?". Some of the important changes in the technology and business landscape are highlighted here:

- New programming languages
  - Go
  - Scala
  - Closure
- New programming paradigms
  - Functional programming
  - Reactive programming
- New frameworks
- New tools
  - Development
  - Code quality
  - Automation testing
  - Deployment
  - Containerizations
- New processes and practices
  - Agile
    - Test-driven development
    - Behavior-driven development
    - Continuous integration
    - Continuous delivery
    - DevOps
  - New devices and opportunities
    - Mobile
    - Cloud

[ 114 ]

Copyrighted material

Page 115 is not part of this book preview.

## Performance

If a web application does not respond within a couple of seconds, there is a very high chance that the user of your application will be disappointed. Performance usually refers to the ability of a system to provide an agreed-upon response time for a defined number of users.

## High resilience

As applications become distributed, the probability of failures increases. How does the application react in the case of localized failures or disruptions? Can it provide basic operations without completely crashing?

This behavior of an application to provide the bare minimum service levels in case of unexpected failures is called resilience.

As more and more applications move towards the Cloud, the resilience of applications becomes important.

We will discuss building highly resilient microservices using *Spring Cloud* and *Spring Data Flow* in Chapter 9, *Spring Cloud* and Chapter 10, *Spring Cloud Data Flow*.

## Scalability

Scalability is a measure of how an application would react when the resources at its disposal are scaled up. If an application supports 10,000 users with a given infrastructure, can it support at least 20,000 users with double the infrastructure?

If a web application does not respond within a couple of seconds, there is a very high chance that the user of your application will be disappointed. Performance usually refers to the ability of a system to provide an agreed-upon response time for a defined number of users.

In the world of Cloud, the scalability of applications becomes even more important. It's difficult to guess how successful a startup might be. Twitter or Facebook might not have expected such success when they were incubated. Their success, for large measure, depends on how they were able to adapt to a multi-fold increase in their user base without affecting the performance.

We will discuss building highly scalable microservices using Spring Cloud and Spring Data Flow in Chapter 9, *Spring Cloud* and Chapter 10, *Spring Cloud Data Flow*.

— [ 116 ] —

Copyrighted material

*Evolution toward Microservices and Cloud-Native Applications*

## Challenges with monolithic applications

Over the last few years, in parallel to working with several small applications, I had the opportunity to work on four different monolithic applications in varied domains—insurance, banking, and health care. All these applications had very similar challenges. In this section, we will start with looking at the characteristics of monoliths and then look at the challenges they bring in.

First of all: What is a monolith? An application with a lot of code—may be greater than 100K lines of code? Yeah.

For me, monoliths are those applications for which getting a release out to production is a big challenge. Applications that fall into this category have a number of user requirements that are immediately needed, but these applications are able to do new feature releases once every few months. Some of these applications even do feature releases once a quarter or sometimes even less as twice a year.

Typically, all monolithic applications have these characteristics:

- **Large size:** Most of these monolithic applications have more than 100K lines of code. Some have codebases with more than a million lines of code.
- **Large teams:** The team size could vary from 20 to 300.
- **Multiple ways of doing the same thing:** Since the team is huge, there is a communication gap. This results in multiple solutions for the same problem in different parts of the application.
- **Lack of automation testing:** Most of these applications have very few unit tests and a complete lack of integration tests. These applications have great dependence on manual testing.

Because of these characteristics, there are a number of challenges faced by these monolithic applications.

## Long release cycles

Making a code change in one part of the monolith may impact some other part of the monolith. Most code changes will need a complete regression cycle. This results in long release cycles.

Because there is a lack of automation testing, these applications depend on manual testing to find defects. Taking the functionality live is a major challenge.

— [ 117 ] —

Copyrighted material

*Evolution toward Microservices and Cloud-Native Applications*

## Difficult to scale

Typically, most monolithic applications are not Cloud-Native, which means that they are not easy to deploy on the Cloud. They depend on manual installation and manual configuration. There is typically a lot of work put in by the operations team before a new application instance is added to the cluster. This makes scaling up and down a big challenge.

The other important challenge is large databases. Typically, monolithic applications have databases running into terabytes (TB). The database becomes the bottleneck when scaling up.

## Adapting new technologies

Most monolithic applications use old technologies. Adding a new technology to the monolith only makes it more complex to maintain. Architects and developers are reluctant to bring in any new technologies.

## Adapting new methodologies

New methodologies such as **agile** need small (four-seven team members), independent teams. The big questions with monolith are these: How do we prevent teams from stepping on each other's toes? How do we create islands that enable teams to work independently? This is a difficult challenge to solve.

## Adapting modern development practices

Modern development practices such as **Test-Driven Development (TDD)**, **Behavior-Driven Development (BDD)** need loosely coupled, testable architecture. If the monolithic application has tightly coupled layers and frameworks, it is difficult to unit test. It makes adapting modern development practices challenging.

## Understanding microservices

The challenges with monolithic applications lead to organizations searching for the silver bullet. How will we be able to make more features live more often?

Pages 119 to 120 are not shown in this preview.

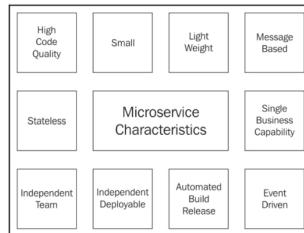
*Evolution toward Microservices and Cloud-Native Applications*

A few important things to note are as follows:

- Modules are identified based on business capabilities. What functionality is the module providing?
- Each module is independently deployable. In the following example, modules 1, 2, and 3 are separate deployable units. If there is a change in the business functionality of module 3, we can individually build and deploy module 3.

## Microservice characteristics

In the previous section, we looked at an example of the microservice architecture. An evaluation of experiences at organizations successful at adapting the microservices architectural style reveals that there are a few characteristics that are shared by teams and architectures. Let's look at some of them:



## Small and lightweight microservices

A good microservice provides a business capability. Ideally, microservices should follow the **single responsibility principle**. Because of this, microservices are generally small in size. Typically, a rule of thumb I use is that it should be possible to build and deploy a microservice within 5 minutes. If the building and deployment takes any longer, it is likely that you are building a larger than recommended microservice.

*Evolution toward Microservices and Cloud-Native Applications*

Some examples of small and lightweight microservices are as follows:

- Product recommendation service
- Email notification service
- Shopping cart service

## Interoperability with message-based communication

The key focus of microservices is on interoperability—communication between systems using diverse technologies. The best way to achieve interoperability is using message-based communication.

## Capability-aligned microservices

It is essential that microservices have a clear boundary. Typically, every microservice has a single identified business capability that it delivers well. Teams have found success adapting the *Bounded Context* concept proposed in the book *Domain-Driven Design* by Eric J. Evans.

Essentially, for large systems, it is very difficult to create one domain model. Evans talks about splitting the system into different bounded contexts. Identifying the right bounded contexts is the key to success with microservice architecture.

## Independently deployable units

Each microservice can be individually built and deployed. In the example discussed earlier, modules 1, 2, and 3 can each be independently built and deployed.

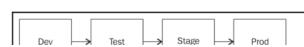
## Stateless

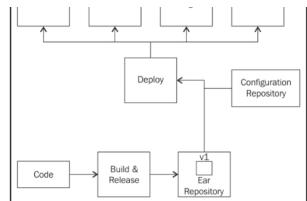
An ideal microservice does not have a state. It does not store any information between requests. All the information needed to create a response is present in the request.

*Evolution toward Microservices and Cloud-Native Applications*

## Automated build and release process

Microservices have automated build and release processes. Take a look at the following figure. It shows a simple build and release process for a microservice:





When a microservice is built and released, a version of the microservice is stored in the repository. The deploy tool has the capability of picking the right version of microservice from the repository, matching it with the configuration needed for the specific environment (from the configuration repository), and deploying the microservice to a specific environment.

Some teams take it a step further and combine the microservice package with the underlying infrastructure needed to run the microservice. The deploy tool will replicate this image and match it with an environment-specific configuration to create an environment.

### Event-driven architecture

Microservices are typically built with event-driven architecture. Let's consider a simple example. Whenever a new customer is registered, there are three things that need to be performed:

- Store customer information to the database
- Mail a welcome kit
- Send an email notification

[ 123 ]

Copyrighted material

#### *Evolution toward Microservices and Cloud-Native Applications*

Let's look at two different approaches to design this.

#### Approach 1 - sequential approach

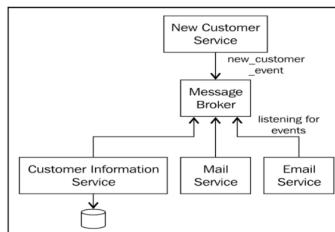
Let's consider three services—`CustomerInformationService`, `MailService`, and `EmailService`, which can provide the capabilities listed earlier. We can create `NewCustomerService` with the following steps:

1. Call `CustomerInformationService` to save customer information to the database.
2. Call `MailService` to mail the welcome kit.
3. Call `EmailService` to send the e-mail notification.

`NewCustomerService` becomes the central place for all business logic. Imagine if we have to do more things when a new customer is created. All that logic would start accumulating and bloating up `NewCustomerService`.

#### Approach 2 - event-driven approach

In this approach, we use a message broker. `NewCustomerService` will create a new event and post it to the message broker. The following figure shows a high-level representation:



[ 124 ]

Copyrighted material

#### *Evolution toward Microservices and Cloud-Native Applications*

The three services—`CustomerInformationService`, `MailService`, and `EmailService`—will be listening on the message broker for new events. When they see the new customer event, they process it and execute the functionality of that specific service.

The key advantage of the event-driven approach is that there is no centralized magnet for all the business logic. Adding a new functionality is easier. We can create a new service to listen for the event on the message broker. Another important thing to note is that we don't need to make changes to any of the existing services.

### Independent teams

The team developing a microservice is typically independent. It contains all the skills needed to develop, test, and deploy a microservice. It is also responsible for supporting the microservice in production.

### Microservice advantages

Microservices have several advantages. They help in keeping up with technology and getting solutions to your customers faster.

### Faster time to market

The faster time to market is one of the key factors in determining the success of an organization.

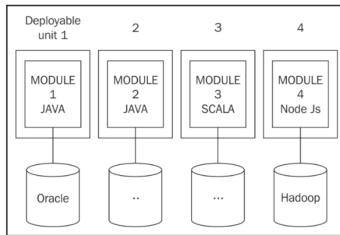
Microservices architecture involves creating small, independently deployable components. Microservice enhancements are easier and less brittle because each microservice focuses on a single business capability. All the steps in the process—building, releasing, deployment, testing, configuration management, and monitoring—are automated. Since the responsibility of a microservice is bounded, it is possible to write great automation unit and integration

tests.

All these factors result in applications being able to react faster to customer needs.

## Technology evolution

There are new languages, frameworks, practices, and automation possibilities emerging every day. It is important that the application architectures allow flexibility to adapt to emerging possibilities. The following figure shows how different services are developed in different technologies:

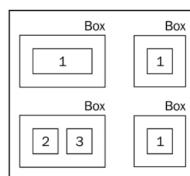


The microservice architecture involves creating small services. Within some boundaries, most organizations give the individual teams the freedom to make some of the technology decisions. This allows teams to experiment with new technologies and innovate faster. This helps applications adapt and stay in tune with the evolution of technology.

## Availability and scaling

The load on different parts of the application is typically very different. For example, in the case of a flight booking application, a customer usually searches multiple times before making a decision on whether to book a flight. The load on the search module would typically be many times more than the load on the booking module. The microservices architecture provides the flexibility of setting up multiple instances of the search service with few instances of the booking service.

The following figure shows how we can scale up specific microservices based on the load:



Microservices 2 and 3 share a single box (the deployment environment). Microservice 1, which has more load, is deployed into multiple boxes.

Another example is the need for start-ups. When a start-up begins its operations, they are typically unaware of the extent to which they might grow. What happens if the demand for an application grows very fast? If they adapt the microservice architecture, it enables them to scale better when the need arises.

## Team dynamics

Development methodologies such as agile advocate small, independent teams. Since microservices are small, it is possible to build small teams around them. Teams are cross-functional, with end-to-end ownership of specific microservices.

Microservice architecture fits in very well with agile and other modern development methodologies.

## Microservice challenges

Microservice architecture has significant advantages. However, there are significant challenges too. Deciding the boundaries of microservices is a challenging but important decision. Since microservices are small, and there would be hundreds of microservices in a large enterprise, having great automation and visibility is critical.

## Increased need for automation

With microservice architecture, you are splitting up a large application into multiple

microservices, so the number of builds, releases, and deployments increases multifold. It would be very inefficient to have manual processes for these steps.

Test automation is critical to enable a faster time to market. Teams should be focused on identifying automation possibilities as they emerge.

## Defining the boundaries of subsystems

Microservices should be intelligent. They are not weak CRUD services. They should model the business capability of the system. They own all the business logic in a bounded context. Having said this, microservices should not be large. Deciding the boundaries of microservices is a challenge. Finding the right boundaries might be difficult on the first go. It is important that as a team gains more knowledge about the business context, the knowledge flows into the architecture and new boundaries are determined. Generally, finding the right boundaries for microservices is an evolutionary process.

A couple of important points to note are as follows:

- Loose coupling and high cohesion are fundamental to any programming and architectural decisions. When a system is loosely coupled, changes in one part should not require a change in other parts.
- Bounded contexts represent autonomous business modules representing specific business capabilities.

*As Sam Newman says in the book Building Microservices—“Specific responsibility enforced by explicit boundaries”. Always think, “What capabilities are we providing to the rest of the domain? ”.*

## Visibility and monitoring

With microservices, one application is split into several microservices. To conquer the complexity associated with multiple microservices and asynchronous event-based collaboration, it is important to have great visibility.

Ensuring high availability means each microservice should be monitored. Automated health management of the microservices becomes important.

[ 128 ]

Copyrighted material

### *Evolution toward Microservices and Cloud-Native Applications*

Debugging problems needs insights into what's happening behind multiple microservices. Centralized logging with aggregation of logs and metrics from different microservices is typically used. Mechanisms such as correlation IDs need to be used to isolate and debug issues.

## Fault tolerance

Let's say we are building a shopping application. What happens if the recommendation microservice is down? How does the application react? Does it completely crash? Or will it let the customer shop? These kinds of situations happen more often as we adapt the microservices architecture.

As we make the services small, the chance that a service is down increases. How the application reacts to these situations becomes an important question. In the earlier example, a fault-tolerant application would show some default recommendations while letting the customer shop.

As we move into microservices architecture, applications should be more fault tolerant. Applications should be able to provide toned-down behavior when services are down.

## Eventual consistency

It is important to have a degree of consistency between microservices in an organization. Consistency between microservices enables similar development, testing, release, deployment, and operational processes across the organization. This enables different developers and testers to be productive when they move across teams. It is important to be not very rigid and have a degree of flexibility within limits so as to not stifle innovation.

### Shared capabilities (enterprise level)

Let's look at a few capabilities that have to be standardized at an enterprise level.

- **Hardware:** What hardware do we use? Do we use Cloud?
- **Code management:** What version control system do we use? What are our practices in branching and committing code?
- **Build and deployment:** How do we build? What tools do we use to automate deployment?
- **Data store:** What kind of data stores do we use?

[ 129 ]

Copyrighted material

Page 130 is not part of this book preview.

### *Evolution toward Microservices and Cloud-Native Applications*

Cloud-Native applications are those that can easily be deployed on the Cloud. These applications share a few common characteristics. We will begin by looking at the Twelve-Factor App—a combination of common patterns among Cloud-Native applications.

## Twelve-Factor App

The Twelve-Factor App evolved from the experiences of engineers at Heroku. It is a list of patterns that are used in Cloud-Native application architectures.

It is important to note that an app here refers to a single deployable unit. Essentially, every microservice is an app (because each microservice is independently deployable).

## Maintain one code base

Each app has one code base in revision control. There can be multiple environments where the app can be deployed. However, all these environments use code from a single codebase. An example antipattern is building a deployable from multiple codebases.

## Dependencies

All dependencies must be explicitly declared and isolated. Typical Java applications use

*Evolution toward Microservices and Cloud-Native Applications*

The following figure shows typical Java applications managing dependencies using Maven:

```

Maven Dependencies
├─ javaee-web-api-7.0.jar - /Users/rangaraokaranam/m2/repo
├─ jstl-1.2.jar - /Users/rangaraokaranam/m2/repo
├─ spring-webmvc-5.0.0.M2.jar - /Users/rangaraokaranam/m2/repo
├─ spring-aop-5.0.0.M2.jar - /Users/rangaraokaranam/m2/repo
├─ spring-beans-5.0.0.M2.jar - /Users/rangaraokaranam/m2/repo
├─ spring-context-5.0.0.M2.jar - /Users/rangaraokaranam/m2/repo
├─ spring-core-5.0.0.M2.jar - /Users/rangaraokaranam/m2/repo
├─ commons-logging-1.2.jar - /Users/rangaraokaranam/m2/repo
├─ spring-expression-5.0.0.M2.jar - /Users/rangaraokaranam/m2/repo
├─ spring-web-5.0.0.M2.jar - /Users/rangaraokaranam/m2/repo
├─ spring-security-web-4.0.1.RELEASE.jar - /Users/rangaraokaranam/m2/repo
├─ aopalliance-1.0.jar - /Users/rangaraokaranam/m2/repo
├─ spring-security-core-4.0.1.RELEASE.jar - /Users/rangaraokaranam/m2/repo
├─ spring-security-config-4.0.1.RELEASE.jar - /Users/rangaraokaranam/m2/repo
├─ hibernate-validator-5.0.2.Final.jar - /Users/rangaraokaranam/m2/repo
├─ validation-api-1.1.0.Final.jar - /Users/rangaraokaranam/m2/repo
├─ jboss-logging-3.1.1.GA.jar - /Users/rangaraokaranam/m2/repo
├─ classmate-1.0.0.jar - /Users/rangaraokaranam/m2/repo
├─ bootstrap-3.3.6.jar - /Users/rangaraokaranam/m2/repo
├─ jquery-1.9.1.jar - /Users/rangaraokaranam/m2/repo
├─ jackson-databind-2.5.3.jar - /Users/rangaraokaranam/m2/repo
├─ jackson-annotations-2.5.0.jar - /Users/rangaraokaranam/m2/repo
├─ jackson-core-2.5.3.jar - /Users/rangaraokaranam/m2/repo
├─ log4j-1.2.17.jar - /Users/rangaraokaranam/m2/repo
├─ junit-4.12.jar - /Users/rangaraokaranam/m2/repo
├─ hamcrest-core-1.3.jar - /Users/rangaraokaranam/m2/repo
├─ spring-test-5.0.0.M2.jar - /Users/rangaraokaranam/m2/repo
├─ mockito-core-1.9.0-rc1.jar - /Users/rangaraokaranam/m2/repo
└─ objenesis-1.0.jar - /Users/rangaraokaranam/m2/repo

```

*Evolution toward Microservices and Cloud-Native Applications*

The following figure shows pom.xml, where the dependencies are managed for a Java application:

```

mastering-spring-chapter-3/pom.xml 33
42      </dependency>
43
44    <dependency>
45      <groupId>org.springframework.security</groupId>
46      <artifactId>spring-security-web</artifactId>
47      <version>4.0.1.RELEASE</version>
48    </dependency>
49
50    <dependency>
51      <groupId>org.springframework.security</groupId>
52      <artifactId>spring-security-config</artifactId>
53      <version>4.0.1.RELEASE</version>
54    </dependency>
55
56    <dependency>
57      <groupId>org.hibernate</groupId>
58      <artifactId>hibernate-validator</artifactId>
59      <version>5.0.2.Final</version>
60    </dependency>
61
62    <dependency>
63      <groupId>org.webjars</groupId>
64      <artifactId>bootstrap</artifactId>
65      <version>3.3.6</version>
66    </dependency>
67
68    <dependency>
69      <groupId>org.webjars</groupId>
70      <artifactId>jquery</artifactId>
71      <version>1.9.1</version>
72    </dependency>
73

```

**Config**

All applications have configuration that varies from one environment to another. Configuration is found at multiple locations; application code, property files, databases, environment variables, JNDI, and system variables are a few examples.

## Disposability

A Twelve-Factor App should promote elastic scaling. Hence, they should be disposable. They can be started and stopped when needed.

A Twelve-Factor App should do the following:

- Have minimum startup time. A long startup time means a long delay before an application can take requests.
- Shut down gracefully.
- Handle hardware failures gracefully.

## Environment parity

All the environments—development, test, staging, and production—should be similar. They should use the same processes and tools. With continuous deployment, they should have very frequently have similar code. This makes finding and fixing problems easier.

## Logs as event streams

Visibility is critical to a Twelve-Factor App. Since applications are deployed on the Cloud and are automatically scaled, it is important that you have a centralized view of what's happening across different instances of the applications.

Treating all logs as stream enables the routing of the log stream to different destinations for viewing and archival purposes. This stream can be used to debug issues, perform analytics, and create alerting systems based on error patterns.

## No distinction of admin processes

Twelve-Factor Apps treat administrative tasks (migrations, scripts) similar to normal application processes.

## Spring projects

As the world moves toward Cloud-Native applications and microservices, Spring projects are not far behind. There are a number of new Spring projects—Spring Boot, Spring Cloud, among others, that solve the problems of the emerging world.

[ 136 ]

Copyrighted material

## Spring Boot

In the era of monoliths, we had the luxury of taking the time to set the frameworks up for an application. However, in the era of microservices, we want to create individual components faster. The Spring Boot project aims to solve this problem.



As the official website highlights, Spring Boot makes it easy to create standalone, production-grade Spring-based applications that you can *just run*. We take an opinionated view of the Spring platform and third-party libraries so that you can get started with minimum fuss.

Spring Boot aims to take an opinionated view—basically making a lot of decisions for us—to developing Spring-based projects.

In the next couple of chapters, we will look at Spring Boot and the different features that enable us to create production-ready applications faster.

## Spring Cloud

Spring Cloud aims to provide solutions to some commonly encountered patterns when building systems on the Cloud:

- **Configuration management:** As we discussed in the Twelve-Factor App section, managing configuration is an important part of developing Cloud-Native applications. Spring Cloud provides a centralized configuration management solution for microservices called Spring Cloud Config.
- **Service discovery:** Service discovery promotes loose coupling between services. Spring Cloud provides integration with popular service discovery options, such as Eureka, ZooKeeper, and Consul.
- **Circuit breakers:** Cloud-Native applications must be fault tolerant. They should be able to handle the failure of backing services gracefully. Circuit breakers play a key role in providing the default minimum service in case of failures. Spring Cloud provides integration with the Netflix Hystrix fault tolerance library.
- **API Gateway:** An API Gateway provides centralized aggregation, routing, and caching services. Spring Cloud provides integration with the API Gateway library Netflix Zuul.

[ 137 ]

Copyrighted material

## Summary

In this chapter, we looked at how the world evolved toward microservices and Cloud-Native applications. We understood how Spring Framework and projects are evolving to meet the needs of today's world with projects such as Spring Boot, Spring Cloud, and Spring Data.

In the next chapter, we will start focusing on Spring Boot. We will look at how Spring Boot makes developing microservices easy.

Pages 139 to 145 are not shown in this preview.

*Building Microservices with Spring Boot*

## Understanding starter projects

Starters are simplified dependency descriptors customized for different purposes. For example, `spring-boot-starter-web` is the starter for building web application, including RESTful, using Spring MVC. It uses Tomcat as the default embedded container. If I want to develop a web application using Spring MVC, all we would need to do is include `spring-boot-starter-web` in our dependencies, and we get the following automatically pre-configured:

- Spring MVC
- Compatible versions of jackson-databind (for binding) and hibernate-validator (for form validation)
- `spring-boot-starter-tomcat` (starter project for Tomcat)

The following code snippet shows some of the dependencies configured in `spring-boot-starter-web`:

```
<dependencies>
    <dependency>
        <groupId>org.springframework.boot</groupId>
        <artifactId>spring-boot-starter</artifactId>
    </dependency>
    <dependency>
        <groupId>org.springframework.boot</groupId>
        <artifactId>spring-boot-starter-tomcat</artifactId>
    </dependency>
    <dependency>
        <groupId>org.hibernate</groupId>
        <artifactId>hibernate-validator</artifactId>
    </dependency>
    <dependency>
        <groupId>com.fasterxml.jackson.core</groupId>
        <artifactId>jackson-databind</artifactId>
    </dependency>
    <dependency>
        <groupId>org.springframework</groupId>
        <artifactId>spring-web</artifactId>
    </dependency>
    <dependency>
        <groupId>org.springframework</groupId>
        <artifactId>spring-webmvc</artifactId>
    </dependency>
</dependencies>
```

*Building Microservices with Spring Boot*

As we can see in the preceding snippet, when we use `spring-boot-starter-web`, we get a lot of frameworks auto-configured.

For the web application we would like to build, we would also want to do some good unit testing and deploy it on Tomcat. The following snippet shows the different starter dependencies that we would need. We would need to add this to our `pom.xml` file:

```
<dependencies>
    <dependency>
        <groupId>org.springframework.boot</groupId>
        <artifactId>spring-boot-starter-web</artifactId>
    </dependency>
    <dependency>
        <groupId>org.springframework.boot</groupId>
        <artifactId>spring-boot-starter-test</artifactId>
        <scope>test</scope>
    </dependency>
    <dependency>
        <groupId>org.springframework.boot</groupId>
        <artifactId>spring-boot-starter-tomcat</artifactId>
        <scope>provided</scope>
    </dependency>
</dependencies>
```

We add three starter projects:

- We've already discussed `spring-boot-starter-web`. It provides us with the frameworks needed to build a web application with Spring MVC.
- The `spring-boot-starter-test` dependency provides the following test frameworks needed for unit testing:
  - **JUnit**: Basic unit test framework
  - **Mockito**: For mocking
  - **Hamcrest, AssertJ**: For readable asserts
  - **Spring Test**: A unit testing framework for spring-context based applications
- The `spring-boot-starter-tomcat` dependency is the default for running web applications. We include it for clarity. `spring-boot-starter-tomcat` is the

We now have our `pom.xml` file configured with the starter parent and the required starter projects. Let's add `spring-boot-maven-plugin` now, which would enable us to run Spring Boot applications.

## Configuring spring-boot-maven-plugin

When we build applications using Spring Boot, there are a couple of situations that are possible:

- We would want to run the applications in place without building a JAR or a WAR
  - We would want to build a JAR and a WAR for later deployment

The `spring-boot-maven-plugin` dependency provides capabilities for both of the preceding situations. The following snippet shows how we can configure `spring-boot-maven-plugin` in an application:

```
<build>
  <plugins>
    <plugin>
      <groupId>org.springframework.boot</groupId>
      <artifactId>spring-boot-maven-plugin</artifactId>
    </plugin>
  </plugins>
</build>
```

The `spring-boot-maven-plugin` dependency provides several goals for a Spring Boot application. The most popular goal is `run` (this can be executed as `mvn spring-boot:run` on the command prompt from the root folder of the project).

## Creating your first Spring Boot launch class

The following class explains how to create a simple Spring Boot launch class. It uses the static run method from the `SpringApplication` class, as shown in the following code snippet:

```
package com.mastering.spring.springboot;
import org.springframework.boot.SpringApplication;
import org.springframework.boot.autoconfigure.SpringBootApplication;
import org.springframework.context.ApplicationContext;
@SpringBootApplication public class Application {
```

```
mvValidator
mvViewResolver
objectNamingStrategy
autoconfigure.AutoConfigurationPackages
autoconfigure.PropertyPlaceholderAutoConfiguration
autoconfigure.condition.BeanTypeRegistry
autoconfigure.context.ConfigurationPropertiesAutoConfiguration
autoconfigure.info.ProjectInfoAutoConfiguration
autoconfigure.internal.CachingMetadataReaderFactory
autoconfigure.jackson.JacksonAutoConfiguration
autoconfigure.jackson.JacksonObjectMapperBuilderConfiguration
autoconfigure.jackson.JacksonAutoConfiguration$JacksonObjectMapperBuilderConfiguration
autoconfigure.jackson.JacksonAutoConfiguration$JacksonObjectMapperConfiguration
autoconfigure.jmx.JmxAutoConfiguration
autoconfigure.web.DispatcherServletAutoConfiguration
autoconfigure.web.DispatcherServletAutoConfiguration$DispatcherServletConfiguration
autoconfigure.web.DispatcherServletAutoConfiguration$DispatcherServletRegistrationConfiguration
autoconfigure.web.EmbeddedServletContainerAutoConfiguration
autoconfigure.web.EmbeddedServletContainerAutoConfiguration$EmbeddedTomcat
autoconfigure.web.ErrorMvcAutoConfiguration
autoconfigure.web.ErrorMvcAutoConfiguration$WhitelabelErrorViewConfiguration
autoconfigure.web.HttpEncodingAutoConfiguration
autoconfigure.web.HttpMessageConvertersAutoConfiguration
autoconfigure.web.HttpMessageConvertersAutoConfiguration$StringHttpMessageConverters
autoconfigure.web.MappingJacksonHttpMessageConvertersConfiguration
autoconfigure.web.MappingJacksonHttpMessageConvertersConfiguration$MappingJackson2HttpMessageConverterConfiguration
autoconfigure.web.MultipartAutoConfiguration
autoconfigure.web.ServerPropertiesAutoConfiguration
autoconfigure.web.WebClientAutoConfiguration
autoconfigure.web.WebClientAutoConfiguration$RestTemplateConfiguration
autoconfigure.web.WebMvcAutoConfiguration
autoconfigure.web.WebMvcAutoConfiguration$EnableWebMvcConfiguration
autoconfigure.web.WebMvcAutoConfiguration$WebMvcAutoConfigurationAdapter
autoconfigure.web.WebMvcAutoConfiguration$WebMvcAutoConfigurationAdapter$FaviconConfiguration
autoconfigure.websocket.WebSocketAutoConfiguration
autoconfigure.websocket.WebSocketAutoConfiguration$TomcatWebSocketConfiguration
context.properties.ConfigurationPropertiesBindingPostProcessor
context.properties.ConfigurationPropertiesBindingPostProcessor.store
```

```
annotation.ConfigurationClassPostProcessor.enhancedConfigurationProcessor  
annotation.ConfigurationClassPostProcessor.importAwareProcessor
```

```

-----
annotation.internalAutowiredAnnotationProcessor
annotation.internalCommonAnnotationProcessor
annotation.internalConfigurationAnnotationProcessor
annotation.internalRequiredAnnotationProcessor
event.internalEventListenerFactory
event.internalEventListenerProcessor
preserveErrorControllerTargetClassPostProcessor
propertySourcesPlaceholderConfigurer
requestContextFilter
requestMappingHandlerAdapter
requestMappingHandlerMapping
resourceHandlerMapping
restTemplateBuilder
serverProperties
simpleControllerHandlerAdapter
spring.http.encoding-autoconfigure.web.HttpEncodingProperties
spring.http.multipart-autoconfigure.web.MultipartProperties
spring.info-autoconfigure.info.ProjectInfoProperties
spring.jackson-objectMapper.jacksonObjectMapperProperties
spring.json-autoconfigure.web.JsonProperties
spring.resources-autoconfigure.web.ResourceProperties
standardJacksonObjectMapperBuilderCustomizer
stringHttpMessageConverter
tomcatEmbeddedServletContainerFactory
viewControllerHandlerMapping
viewResolver
websocketContainerCustomizer

```

Important things to think about are as follows:

- Where are these beans defined?
- How are these beans created?

That's the magic of Spring auto-configuration.

Whenever we add a new dependency to a Spring Boot project, Spring Boot auto-configuration automatically tries to configure the beans based on the dependency.

For example, when we add a dependency in `spring-boot-starter-web`, the following beans are auto-configured:

- `basicErrorHandler`, `handlerExceptionResolver`: Basic exception handling. Shows a default error page when an exception occurs.
- `beanNameHandlerMapping`: Used to resolve paths to a handler (controller).

----- [ 155 ] -----

Copyrighted material

#### *Building Microservices with Spring Boot*

- `characterEncodingFilter`: Provides default character encoding UTF-8.
- `dispatcherServlet`: `DispatcherServlet` is the Front Controller in Spring MVC applications.
- `jacksonObjectMapper`: Translates objects to JSON and JSON to objects in REST services.
- `messageConverters`: The default message converters to convert from objects into XML or JSON and vice versa.
- `multipartResolver`: Provides support to upload files in web applications.
- `mvcValidator`: Supports validation of HTTP requests.
- `viewResolver`: Resolves a logical view name to a physical view.
- `propertySourcesPlaceholderConfigurer`: Supports the externalization of application configuration.
- `requestContextFilter`: Defaults the filter for requests.
- `restTemplateBuilder`: Used to make calls to REST services.
- `tomcatEmbeddedServletContainerFactory`: Tomcat is the default embedded servlet container for Spring Boot-based web applications.

In the next section, let's look at some of the starter projects and the auto-configuration they provide.

## Starter projects

The following table shows some of the important starter projects provided by Spring Boot:

Starter	Description
<code>spring-boot-starter-web-services</code>	This is a starter project to develop XML-based web services.
<code>spring-boot-starter-web</code>	This is a starter project to build Spring MVC-based web applications or RESTful applications. It uses Tomcat as the default embedded servlet container.
<code>spring-boot-starter-activemq</code>	This supports message-based communication using JMS on ActiveMQ.
<code>spring-boot-starter-integration</code>	This supports the Spring Integration Framework that provides implementations for Enterprise Integration Patterns.

----- [ 156 ] -----

Copyrighted material

Pages 157 to 161 are not shown in this preview.

#### *Building Microservices with Spring Boot*

```

@Test
public void welcome() throws Exception {
    mockMvc.perform(MockMvcRequestBuilders.get("/welcome")
        .accept(MediaType.APPLICATION_JSON))
        .andExpect(status().isOk())
        .andExpect(content().string(
            equalTo("Hello World")));
}
}

```

In the preceding unit test, we will launch up a Mock MVC instance with `BasicController`. A few quick things to note are as follows:

- `@RunWith(SpringRunner.class)`: `SpringRunner` is a shortcut to the `SpringJUnit4ClassRunner` annotation. This launches up a simple Spring context for unit testing.
- `@WebMvcTest(BasicController.class)`: This annotation can be used along with `SpringRunner` to write simple tests for Spring MVC controllers. This will load only the beans annotated with Spring-MVC-related annotations. In this example we are launching a Web MVC Test context with the class under test being `BasicController`.
- `@Autowired private MockMvc mvc`: Autowires the `MockMvc` bean that can be

used to make requests.

- `mvc.perform(MockMvcRequestBuilders.get("/welcome").accept(MediaType.APPLICATION_JSON))`: Performs a request to /welcome with the Accept header value application/json.
- `andExpect(status().isOk())`: Expects that the status of the response is 200 (success).
- `andExpect(content().string(equalTo("Hello World")))`: Expects that the content of the response is equal to "Hello World".

## Integration testing

When we do integration testing, we would want to launch the embedded server with all the controllers and beans that are configured. This code snippet shows how we can create a simple integration test:

```
@RunWith(SpringRunner.class)
@SpringBootTest(classes = Application.class,
    webEnvironment = SpringBootTest.WebEnvironment.RANDOM_PORT)
public class BasicControllerIT {
```

[ 162 ]

Copyrighted material

### *Building Microservices with Spring Boot*

```
private static final String LOCAL_HOST =
"http://localhost:";

@LocalServerPort
private int port;

private TestRestTemplate template = new TestRestTemplate();

@Test
public void welcome() throws Exception {
    ResponseEntity<String> response = template
        .getForEntity(createURL("/welcome"), String.class);
    assertThat(response.getBody(), equalTo("Hello World"));
}

private String createURL(String uri) {
    return LOCAL_HOST + port + uri;
}
```

A few important things to note are as follows:

- `@SpringBootTest(classes = Application.class, webEnvironment = SpringBootTest.WebEnvironment.RANDOM_PORT)`: Provides additional functionality on top of the Spring TestContext. Provides support to configure the port for fully running the container and TestRestTemplate (to execute requests).
- `@LocalServerPort private int port:SpringBootTest` would ensure that the port on which the container is running is autowired into the port variable.
- `private String createURL(String uri):` The method to append the local host URL and port to the URI to create a full URL.
- `private TestRestTemplate template = new TestRestTemplate():` TestRestTemplate is typically used in integration tests. It provides additional functionality on top of RestTemplate, which is especially useful in the integration test context. It does not follow redirects so that we can assert response location.
- `template.getForEntity(createURL("/welcome"), String.class):` Executes a get request for the given URL.
- `assertThat(response.getBody(), equalTo("Hello World")):` Asserts that the response body content is "Hello World".

[ 163 ]

Copyrighted material

### *Building Microservices with Spring Boot*

## Simple REST method returning an object

In the previous method, we returned a string. Let's create a method that returns a proper JSON response. Take a look at the following method:

```
@GetMapping("/welcome-with-object")
public WelcomeBean welcomeWithObject() {
    return new WelcomeBean("Hello World");
}
```

This preceding method returns a simple `WelcomeBean` initialized with a message: "Hello World".

## Executing a request

Let's send a test request and see what response we get. The following screenshot shows the output:



The response for the `http://localhost:8080/welcome-with-object` URL is shown as follows:

```
{"message":"Hello World"}
```

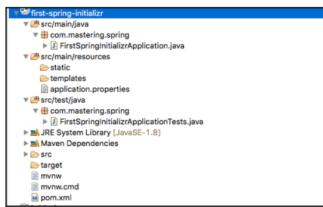
The question that needs to be answered is this: how does the `WelcomeBean` object that we returned get converted into JSON?

Again, it's the magic of Spring Boot auto-configuration. If Jackson is on the classpath of an application, instances of the default object to JSON (and vice versa) converters are auto-configured by Spring Boot.

[ 164 ]

Copyrighted material

This will import the project into Eclipse. The following screenshot shows the structure of the project in Eclipse:



Let's look at some of the important files from the generated project.

### pom.xml

The following snippet shows the dependencies that are declared:

```
<dependencies>
    <dependency> <groupId>org.springframework.boot</groupId>
        <artifactId>spring-boot-starter-web</artifactId> </dependency>
    <dependency> org.springframework.boot </groupId>
        <artifactId>spring-boot-starter-data-jpa</artifactId> </dependency>
    <dependency> org.springframework.boot </groupId>
        <artifactId>spring-boot-starter-actuator</artifactId> </dependency>
    <dependency> org.springframework.boot </groupId>
        <artifactId>spring-boot-devtools</artifactId> </dependency>
    <dependency> org.springframework.boot </groupId>
        <artifactId>spring-boot-starter-test</artifactId> </dependency>
</dependencies>
```

A few other important observations are as follows:

- The packaging for this component is .jar
- org.springframework.boot:spring-boot-starter-parent is declared as the parent POM
- <java.version>1.8</java.version>: The Java version is 1.8
- Spring Boot Maven Plugin (org.springframework.boot:spring-boot-maven-plugin) is configured as a plugin

[ 185 ]

Copyrighted material

### FirstSpringInitializrApplication.java class

FirstSpringInitializrApplication.java is the launcher for Spring Boot:

```
package com.mastering.spring;
import org.springframework.boot.SpringApplication;
import org.springframework.boot.autoconfigure.SpringBootApplication;
@SpringBootApplication
public class FirstSpringInitializrApplication {
    public static void main(String[] args) {
        SpringApplication.run(FirstSpringInitializrApplication.class,
            args);
    }
}
```

### FirstSpringInitializrApplicationTests class

FirstSpringInitializrApplicationTests contains the basic context that can be used to start writing the tests as we start developing the application:

```
package com.mastering.spring;
import org.junit.Test;
import org.junit.runner.RunWith;
import org.springframework.boot.test.context.SpringBootTest;
import org.springframework.test.context.junit4.SpringRunner;

@RunWith(SpringRunner.class)
@SpringBootTest
public class FirstSpringInitializrApplicationTests {

    @Test
    public void contextLoads() {
    }
}
```

### A quick peek into auto-configuration

Auto-configuration is one of the most important features of Spring Boot. In this section, we will take a quick peek behind the scenes to understand how Spring Boot auto-configuration works.

[ 186 ]

Copyrighted material

Most of the Spring Boot auto-configuration magic comes from spring-boot-autoconfigure-{version}.jar. When we start any Spring Boot applications, a number of beans get auto-configured. How does this happen?

The following screenshot shows an extract from spring.factories from spring-boot-autoconfigure-{version}.jar. We have filtered out some of the configuration in the interest of space:

```
spring.factories
31 org.springframework.boot.autoconfigure.data.jpa.JpaRepositories$EnableAutoConfiguration,
32 org.springframework.boot.autoconfigure.data.mongo.MongoDatabaseConfiguration,
33 org.springframework.boot.autoconfigure.data.neo4j.Neo4jDatabaseConfiguration,
34 org.springframework.boot.autoconfigure.data.neo4j.Neo4jAutoConfiguration,
35 org.springframework.boot.autoconfigure.data.neo4j.Neo4jRepositories$EnableAutoConfiguration,
```

```

36 org.springframework.boot.autoconfigure.data.web.SpringDataWebAutoConfiguration,
37 org.springframework.boot.autoconfigure.data.redis.RedisAutoConfiguration,
38 org.springframework.boot.autoconfigure.data.rdbms.RelationalAutoConfiguration,
39 org.springframework.boot.autoconfigure.data.redis.RedislettuceAutoConfiguration,
40 org.springframework.boot.autoconfigure.data.web.SpringDataWebAutoConfiguration,
41 org.springframework.boot.autoconfigure.elasticsearch.JestAutoConfiguration,
42 org.springframework.boot.autoconfigure.groovytemplate.GroovyTemplateAutoConfiguration,
43 org.springframework.boot.autoconfigure.gson.GsonAutoConfiguration,
44 org.springframework.boot.autoconfigure.h2.H2ConsoleAutoConfiguration,
45 org.springframework.boot.autoconfigure.hateoas.HypermediaAutoConfiguration,
46 org.springframework.boot.autoconfigure.hazelcast.HazelcastAutoConfiguration,
47 org.springframework.boot.autoconfigure.jdbc.DataSourceAutoConfiguration,
48 org.springframework.boot.autoconfigure.info.ProjectInfoAutoConfiguration,
49 org.springframework.boot.autoconfigure.integration.IntegrationAutoConfiguration,
50 org.springframework.boot.autoconfigure.jackson.JacksonAutoConfiguration,
51 org.springframework.boot.autoconfigure.jdbc.DataSourceAutoConfiguration,
52 org.springframework.boot.autoconfigure.jdbc.EmbeddedDatabaseAutoConfiguration,
53 org.springframework.boot.autoconfigure.jdbc.JdbcDataSourceAutoConfiguration,
54 org.springframework.boot.autoconfigure.jdbc.XADataSourceAutoConfiguration,
55 org.springframework.boot.autoconfigure.jdbc.JdbcTransactionManagerAutoConfiguration,
56 org.springframework.boot.autoconfigure.jms.JmsAutoConfiguration,
57 org.springframework.boot.autoconfigure.jms.JmsAutoConfiguration,
58 org.springframework.boot.autoconfigure.jms.JndiConnectionFactoryAutoConfiguration,
59 org.springframework.boot.autoconfigure.jms.ActiveMQJmsAutoConfiguration,
60 org.springframework.boot.autoconfigure.jms.ArtemisJmsAutoConfiguration,
61 org.springframework.boot.autoconfigure.jms.HornetqJmsAutoConfiguration,
62 org.springframework.boot.autoconfigure.flyway.FlywayAutoConfiguration,

```

The preceding list of auto-configuration classes is run whenever a Spring Boot application is launched. Let's take a quick look at one of them:

```
org.springframework.boot.autoconfigure.web.WebMvcAutoConfiguration.
```

[ 187 ]

Copyrighted material

Pages 188 to 192 are not shown in this preview.

#### *Extending Microservices*

As we can see in the preceding two examples, Spring Boot provides good default exception handling. In the next section, we will focus on understanding how the application reacts to custom exceptions.

### Throwing a custom exception

Let's create a custom exception and throw it from a service. Take a look at the following code:

```

public class TodoNotFoundException extends RuntimeException {
    public TodoNotFoundException(String msg) {
        super(msg);
    }
}

```

It's a very simple piece of code that defines `TodoNotFoundException`.

Now let's enhance our `TodoController` class to throw `TodoNotFoundException` when a `todo` with a given ID is not found:

```

@GetMapping(path = "/users/{name}/todos/{id}")
public Todo retrieveTodo(@PathVariable String name,
@PathVariable int id) {
    Todo todo = todoService.retrieveTodo(id);
    if (todo == null) {
        throw new TodoNotFoundException("Todo Not Found");
    }
    return todo;
}

```

If `todoService` returns a null `todo`, we throw `TodoNotFoundException`.

When we execute the service with a GET request to a nonexistent `todo` (`http://localhost:8080/users/Jack/todos/222`), we get the response shown in the following code snippet:

```

{
    "timestamp": 1484029048788,
    "status": 500,
    "error": "Internal Server Error",
    "exception",
    "com.mastering.spring.springboot.bean.TodoNotFoundException",
    "message": "Todo Not Found",
    "path": "/users/Jack/todos/222"
}

```

[ 193 ]

Copyrighted material

#### *Extending Microservices*

As we can see, a clear exception response is sent back to the service consumer. However, there is one thing that can be improved further—the response status. When a resource is not found, it is recommended that you return a 404 – Resource Not Found status. We will look at how to customize the response status in the next example.

### Customizing the exception message

Let's look at how to customize the preceding exception and return the proper response status with a customized message.

Let's create a bean to define the structure of our custom exception message:

```

public class ErrorResponse {
    private Date timestamp = new Date();
    private String message;
    private String details;

    public ErrorResponse(String message, String details) {
        super();
        this.message = message;
        this.details = details;
    }

    public Date getTimestamp() {
        return timestamp;
    }

    public String getMessage() {
        return message;
    }

    public String getDetails() {
        return details;
    }
}

```

We have created a simple exception response bean with an auto-populated timestamp with

*Extending Microservices*

When `TodoNotFoundException` is thrown, we would want to return a response using the `ErrorResponse` bean. The following code shows how we can create a global exception handling for `TodoNotFoundException.class`:

```
@ControllerAdvice
@RestController
public class RestResponseEntityExceptionHandler
    extends ResponseEntityExceptionHandler
{
    @ExceptionHandler(TodoNotFoundException.class)
    public final ResponseEntity<ErrorResponse>
        todoNotFound(TodoNotFoundException ex) {
        ErrorResponse exceptionResponse =
            new ErrorResponse( ex.getMessage(),
                "Any details you would want to add");
        return new ResponseEntity<ErrorResponse>
            (exceptionResponse, new HttpHeaders(),
                HttpStatus.NOT_FOUND);
    }
}
```

Some important things to note are as follows:

- `RestResponseEntityExceptionHandler` extends `ResponseEntityExceptionHandler`: We are extending `ResponseEntityExceptionHandler`, which is the base class provided by Spring MVC for centralised exception handling `ControllerAdvice` classes.
- `@ExceptionHandler(TodoNotFoundException.class)`: This defines that the method to follow will handle the specific exception `TodoNotFoundException.class`. Any other exceptions for which custom exception handling is not defined will follow the default exception handling provided by Spring Boot.
- `ErrorResponse exceptionResponse = new ErrorResponse(ex.getMessage(), "Any details you would want to add");`: This creates a custom exception response.
- `new ResponseEntity<ErrorResponse>(exceptionResponse, new HttpHeaders(), HttpStatus.NOT_FOUND)`: This is the definition to return a 404 Resource Not Found response with the custom exception defined earlier.

Pages 196 to 197 are not shown in this preview.

*Extending Microservices*

A REST API must not define fixed resource names or hierarchies (an obvious coupling of client and server). Servers must have the freedom to control their own namespace. Instead, allow servers to instruct clients on how to construct appropriate URIs, such as is done in HTML forms and URI templates, by defining those instructions within media types and link relations. A REST API should be entered with no prior knowledge beyond the initial URI (bookmark) and set of standardized media types that are appropriate for the intended audience (i.e., expected to be understood by any client that might use the API). From that point on, all application state transitions must be driven by client selection of server-provided choices that are present in the received representations or implied by the user's manipulation of those representations. The transitions may be determined (or limited by) the client's knowledge of media types and resource communication mechanisms, both of which may be improved on-the-fly (e.g., code-on-demand).

An example response with HATEOAS links is shown here. This is the response to the `/todos` request in order to retrieve all todos:

```
{
    "_embedded" : {
        "todos" : [ {
            "user" : "Jill",
            "desc" : "Learn Hibernate",
            "done" : false,
            "_links" : {
                "self" : {
                    "href" : "http://localhost:8080/todos/1"
                },
                "todo" : {
                    "href" : "http://localhost:8080/todos/1"
                }
            }
        } ]
    },
    "_links" : {
        "self" : {
            "href" : "http://localhost:8080/todos"
        },
        "profile" : {
            "href" : "http://localhost:8080/profile/todos"
        },
        "search" : {
            "href" : "http://localhost:8080/todos/search"
        }
    }
}
```

*Extending Microservices*

The preceding response includes links to the following:

- Specific todos (<http://localhost:8080/todos/1>)
- Search resource (<http://localhost:8080/todos/search>)

If the service consumer wants to do a search, it has the option of taking the search URL from the response and sending the search request to it. This would reduce coupling between the service provider and the service consumer.

## Sending HATEOAS links in response

Now that we understand what HATEOAS is, let's look at how we can send links related to a resource in the response.

### Spring Boot starter HATEOAS

Spring Boot has a specific starter for HATEOAS called `spring-boot-starter-hateoas`. We need to add it to the `pom.xml` file.

The following code snippet shows the dependency block:

```
<dependency>
  <groupId>org.springframework.boot</groupId>
  <artifactId>spring-boot-starter-hateoas</artifactId>
</dependency>
```

One of the important dependencies of `spring-boot-starter-hateoas` is `spring-hateoas`, which provides the HATEOAS features:

```
<dependency>
  <groupId>org.springframework.hateoas</groupId>
  <artifactId>spring-hateoas</artifactId>
</dependency>
```

Let's enhance the `retrieveTodo` resource (`/users/{name}/todos/{id}`) to return a link to retrieve all todos (`/users/{name}/todos`) in the response:

```
@GetMapping(path = "/users/{name}/todos/{id}")
public Resource<Todo> retrieveTodo(
    @PathVariable String name, @PathVariable int id) {
    Todo todo = todoService.retrieveTodo(id);
    if (todo == null) {
        throw new TodoNotFoundException("Todo Not Found");
    }
}
```

[ 199 ]

Copyrighted material

### Extending Microservices

```

}
Resource<Todo> todoResource = new Resource<Todo>(todo);
ControllerLinkBuilder linkTo =
linkTo(methodOn(this.getClass()).retrieveTodos(name));
todoResource.add(linkTo.withRel("parent"));
return todoResource;
}
```

Some important points to note are as follows:

- `ControllerLinkBuilder linkTo = linkTo(methodOn(this.getClass()).retrieveTodos(name))`: We want to get a link to the `retrieveTodos` method on the current class
- `linkTo.withRel("parent")`: Relationship with the current resource is parent

The following snippet shows the response when a GET request is sent to `http://localhost:8080/users/Jack/todos/1`:

```
{
    "id": 1,
    "user": "Jack",
    "desc": "Learn Spring MVC",
    "targetDate": 1484038262110,
    "done": false,
    "_links": {
        "parent": {
            "href": "http://localhost:8080/users/Jack/todos"
        }
    }
}
```

The `_links` section will contain all the links. Currently, we have one link with the relation `parent` and `href` as `http://localhost:8080/users/Jack/todos`.

 If you have problems executing the preceding request, try executing using an Accept header—`application/json`.

HATEOAS is not something that is commonly used in most of the resources today. However, it has the potential to be really useful in reducing the coupling between the service provider and the consumer.

[ 200 ]

Copyrighted material

### Extending Microservices

## Validation

A good service always validates data before processing it. In this section, we will look at the Bean Validation API and use its reference implementation to implement validation in our services.

The Bean Validation API provides a number of annotations that can be used to validate beans. The JSR 349 specification defines Bean Validation API 1.1. Hibernate-validator is the reference implementation. Both are already defined as dependencies in the `spring-boot-web-starter` project:

- `hibernate-validator-5.2.4.Final.jar`
- `validation-api-1.1.0.Final.jar`

We will create a simple validation for the `createTodo` service method.

Creating validations involves two steps:

1. Enabling validation on the controller method.
2. Adding validations on the bean.

### Enabling validation on the controller method

It's very simple to enable validation on the controller method. The following snippet shows

an example:

```
@RequestMapping(method = RequestMethod.POST,  
path = "/users/{name}/todos")  
ResponseEntity<?> add(@PathVariable String name  
@Valid @RequestBody Todo todo) {
```

The `@Valid`(package `javax.validation`) annotation is used to mark a parameter for validation. Any validation that is defined in the `Todo` bean is executed before the `add` method is executed.

[ 201 ]

Copyrighted material

Pages 202 to 208 are not shown in this preview.

### Extending Microservices

The following screenshot shows the list of controller-exposing services. When we click on any controller, it expands to show the list of request methods and URLs each controller supports:

The screenshot shows the Swagger UI interface. At the top, there's a navigation bar with tabs like 'Swagger', 'Edit', 'Key', and 'Explore'. Below that is a header for 'Apache 2.0' and three controller sections: 'basic-controller', 'basic-error-controller', and 'todo-controller'. The 'todo-controller' section is expanded, showing a table of operations:

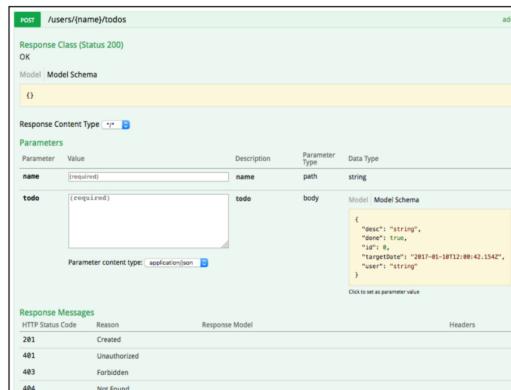
Method	URL	Description
GET	/users/dummy-service	
GET	/users/{name}/todos	retrieveTodos
POST	/users/{name}/todos	add
GET	/users/{name}/todos/{id}	retrieveTodo

[ 209 ]

Copyrighted material

### Extending Microservices

The following screenshot shows the details for the POST service to create a `todo` for the user in Swagger UI:



The screenshot shows the Swagger UI for the `POST /users/{name}/todos` endpoint. The 'Parameters' section shows two parameters: `name` (required, path) and `todo` (required, body). The 'Response Content Type' dropdown is set to `application/json`. To the right, the 'Model Schema' for the `todo` body is displayed as a JSON object:

```
{  
  "done": "string",  
  "done": true,  
  "targetDate": "2017-01-10T12:00:42.154Z",  
  "user": "string"  
}
```

The 'Response Messages' section lists HTTP status codes and their corresponding reasons and headers:

HTTP Status Code	Reason	Response Model	Headers
201	Created		
401	Unauthorized		
403	Forbidden		
404	Not Found		

Some important things to note are as follows:

- **Parameters** show all the important parameters including the request body
- The **Parameter Type** body (for the `todo` parameter) shows the expected structure for the body of the request
- The **Response Messages** sections show different HTTP status codes returned by the service

The Swagger UI provides an excellent way to expose service definitions for your API without a lot of additional effort.

[ 210 ]

## Customizing Swagger documentation using annotations

The Swagger UI also provides annotations to further customize your documentation.

Listed here is some of the documentation to retrieve the todos service:

```
"/users/{name}/todos": {
  "get": {
    "tags": [
      "todo-controller"
    ],
    "summary": "retrieveTodos",
    "operationId": "retrieveTodosUsingGET",
    "consumes": [
      "application/json"
    ],
    "produces": [
      "application/json"
    ]
  }
}
```

As you can see, the documentation generated is very raw. There are a number of things we can improve in the documentation to describe the services better. Here are a couple of examples:

- Provide a better summary
- Add application/JSON to produces

Swagger provides annotations we can add to our RESTful services in order to customize the documentation. Let's add a few annotations to the controller in order to improve the documentation:

```
@ApiOperation(
  value = "Retrieve all todos for a user by passing in his name",
  notes = "A list of matching todos is returned. Current pagination is not supported.",
  response = Todo.class,
  responseContainer = "List",
  produces = "application/json")
@GetMapping("/users/{name}/todos")
public List<Todo> retrieveTodos(@PathVariable String name) {
  return todoService.retrieveTodos(name);
}
```

[ 211 ]

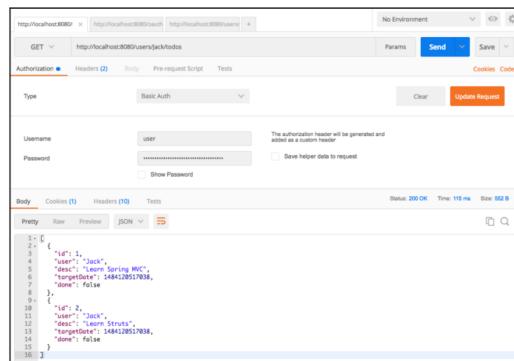
Pages 212 to 214 are not shown in this preview.

An example is shown in the following code snippet:

```
2017-01-11 13:11:58.696 INFO 3888 --- [ restartedMain]
b.a.s.AuthenticationManagerConfiguration :
Using default security password: 3fb5564a-ce53-413b-9911-8ade17b2f478
2017-01-11 13:11:58.771 INFO 3888 --- [ restartedMain]
o.s.s.web.DefaultSecurityFilterChain : Creating filter chain: Ant
[pattern="/css/**"], []
```

Underlined in the preceding code snippet is the default security password printed in the log.

We can use Postman to fire a request with basic authentication. The following screenshot shows how basic authentication details can be sent along with a request:



As you can see, authentication succeeds and we get a proper response back.

[ 215 ]

We can configure the user ID and password of our choice in `application.properties`, as shown here:

```
security.user.name=user-name
security.user.password=user-password
```

Spring Security also provides options to authenticate with LDAP or JDBC or any other data source with user credentials.

## Integration testing

The integration test we wrote for the service earlier will start failing because of invalid

credentials. We will now update the integration test to supply basic authentication credentials:

```
private TestRestTemplate template = new TestRestTemplate();
HttpHeaders headers = createHeaders("user-name", "user-password");

HttpHeaders createHeaders(String username, String password) {
    HttpHeaders headers = new HttpHeaders();
    String auth = username + ":" + password;
    byte[] encodedAuth = Base64.getEncoder().encode(
        auth.getBytes(Charset.forName("US-ASCII")));
    String authHeader = "Basic " + new String(encodedAuth);
    headers.set("Authorization", authHeader);
}
}

@Test
public void retrieveTodos() throws Exception {
    String expected = "[{" +
        "(id:1,user:Jack,desc:'Learn Spring MVC',done:false)" + "," +
        "(id:2,user:Jack,desc:'Learn Struts',done:false)" + "}";
    ResponseEntity<String> response = template.exchange(
        createUrl("/users/Jack/todos"), HttpMethod.GET,
        new HttpEntity<String>(null, headers),
        String.class);
    JSONAssert.assertEquals(expected, response.getBody(), false);
}
```

Some important things to note are as follows:

- `createHeaders("user-name", "user-password")`: This method creates `Base64.getEncoder().encode` basic authentication headers

[ 216 ]

Copyrighted material

## Extending Microservices

```
• ResponseEntity<String> response = template.exchange(createUrl("/users/Jack/todos"),
    HttpMethod.GET, new HttpEntity<String>(null, headers),
    String.class): The key change is the use of HttpEntity to supply the headers that we created earlier to the REST template
```

## Unit testing

We would not want to use security for our unit tests. The following code snippet shows how we can disable security for the unit test:

```
@RunWith(SpringRunner.class)
@WebMvcTest(value = TodoController.class, secure = false)
public class TodoControllerTest {
```

The key part is the `secure = false` parameter on the `WebMvcTest` annotation. This will disable Spring Security for the unit test.

## OAuth 2 authentication

OAuth is a protocol that provides flows in order to exchange authorization and authentication information between a range of web-enabled applications and services. It enables third-party applications to get restricted access to user information from a service, for example, Facebook, Twitter, or GitHub.

Before we get into the details, it would be useful to review the terminology typically used with respect to OAuth 2 authentication.

Let's consider an example. Let's say we want to expose the Todo API to third-party applications on the internet.

The following are the important players in a typical OAuth 2 exchange:

- **Resource owner**: This is the user of the third-party application that wants to use our Todo API. It decides how much of the information available with our API can be made available to the third-party application.
- **Resource server**: This hosts the Todo API, the resource we want to secure.
- **Client**: This is the third-party application that wants to consume our API.
- **Authorization server**: This is the server that provides the OAuth service.

[ 217 ]

Copyrighted material

Pages 218 to 225 are not shown in this preview.

## Extending Microservices

## Caching

Caching data from services plays a crucial role in improving the performance and scalability of applications. In this section, we will look at the implementation options that Spring Boot provides.

Spring provides a caching abstraction based on annotations. We will start with using Spring caching annotations. Later, we will introduce `JSR-107` caching annotations and compare them with Spring abstractions.

## Spring-boot-starter-cache

Spring Boot provides a starter project for caching `spring-boot-starter-cache`. Adding this to an application brings in all the dependencies to enable `JSR-107` and Spring caching annotations. The following code snippet shows the dependency details for `spring-boot-starter-cache`. Let's add this to our file `pom.xml`:

```
<dependency>
    <groupId>org.springframework.boot</groupId>
    <artifactId>spring-boot-starter-cache</artifactId>
</dependency>
```

## Enabling caching

Before we can start using caching, we need to enable caching on the application. The following code snippet shows how we can enable caching:

```
@EnableCaching
@SpringBootApplication
public class Application {
```

`@EnableCaching` would enable caching in a Spring Boot application.

## Caching data

Now that we have enabled caching, we can add the `@Cacheable` annotation to the methods where we want to cache the data. The following code snippet shows how to enable caching on `retrieveTodos`:

```
@Cacheable("todos")
public List<Todo> retrieveTodos(String user) {
```

In the preceding example, the `todos` for a specific user are cached. On the first call to the method for a specific user, the `todos` will be retrieved from the service. On subsequent calls for the same user, the data will be returned from the cache.

Spring also provides conditional caching. In the following snippet, caching is enabled only if the specified condition is satisfied:

```
@Cacheable(cacheNames="todos", condition="#user.length < 10")
public List<Todo> retrieveTodos(String user) {
```

Spring also provides additional annotations to evict data from the cache and add some custom data to cache. A few important ones are listed as follows:

- `@CachePut`: Used to explicitly add data to the cache
- `@CacheEvict`: Used to remove stale data from the cache
- `@Caching`: Allows multiple nested `@Cacheable`, `@CachePut`, and `@CacheEvict` annotations to be used on the same method

## JSR-107 caching annotations

*JSR-107* aims to standardize caching annotations. Listed here are some of the important *JSR-107* annotations:

- `@CacheResult`: Similar to `@Cacheable`
- `@CacheRemove`: Similar to `@CacheEvict`; `@CacheRemove` supports conditional eviction if an exception occurs
- `@CacheRemoveAll`: Similar to `@CacheEvict(allEntries=true)`; used to remove all entries from the cache

*JSR-107* and Spring's caching annotations are fairly similar in terms of the features they offer. Either of them is a good choice. We lean slightly toward *JSR-107* because it's a standard. However, make sure you are not using both in the same project.

## Auto-detection order

When caching is enabled, Spring Boot auto-configuration starts looking for a caching provider. The following list shows the order in which Spring Boot searches for caching providers. The list is in order of decreasing preference:

- JCache (*JSR-107*) (EhCache 3, Hazelcast, Infinispan, and so on)
- EhCache 2.x
- Hazelcast
- Infinispan
- Couchbase
- Redis
- Caffeine
- Guava
- Simple

## Summary

Spring Boot makes developing Spring-based applications easy. It enables us to create production-ready applications very quickly.

In this chapter, we covered how to add features such as exception handling, caching, and internationalization to our application. We discussed the best practices of documenting REST services using Swagger. We looked at the basics of securing our microservice with Spring Security.

In the next chapter, we will shift our attention toward advanced features in Spring Boot. We will look at how to provide monitoring on top of our REST services, learn how to deploy the microservice to the Cloud, and understand how to become more productive when developing applications with Spring Boot.



## Advanced Spring Boot Features

In the previous chapter, we extended our microservice with exception handling, HATEOAS, caching, and internationalization. In this chapter, let's turn our attention to deploying our services to production. To be able to deploy the services to production, we need to be able to set up and create functionality to configure, deploy, and monitor services.

The following are some of the questions we will answer during this chapter:

- How to externalize application configuration?
- How to use profiles to configure environment-specific values?
- How to deploy our application to the Cloud?
- What is an embedded server? How can you use Tomcat, Jetty, and Undertow?
- What monitoring features does Spring Boot Actuator provide?
- How can you be a more productive developer with Spring Boot?

Copyrighted material

Pages 230 to 236 are not shown in this preview.

### *Advanced Spring Boot Features*

```

        }
        public String getService1Url() {
            return service1Url;
        }
        public void setService1Url(String service1Url) {
            this.service1Url = service1Url;
        }
        public int getService1Timeout() {
            return service1Timeout;
        }
        public void setService1Timeout(int service1Timeout) {
            this.service1Timeout = service1Timeout;
        }
    }
}

```

A couple of important things to note are as follows:

- `@ConfigurationProperties("application")` is the annotation for an externalized configuration. We can add this annotation to any class to bind to external properties. The value in the double quotes—`application`—is used as a prefix while binding external configuration to this bean.
- We are defining multiple configurable values in the bean.
- Getters and setters are needed since binding happens through Java beans property descriptors.

The following snippet shows how the values for these properties can be defined in `application.properties`:

```
application.enableSwitchForService1=true
application.service1Url=http://abc-dev.service.com/somethingelse
application.service1timeout=250
```

A couple of important things to note are as follows:

- `application`: The prefix is defined as part of `@ConfigurationProperties("application")` while defining the configuration bean
- Values are defined by appending the prefix to the name of the property

[ 237 ]

Copyrighted material

Page 238 is not part of this book preview.

### *Advanced Spring Boot Features*

```
required type 'int' for property 'service1Timeout'; nested exception is
org.springframework.core.convert.ConverterNotFoundException: No converter
found capable of converting from type [java.lang.String] to type [int]
```

```
Action:
Update your application's configuration
```

## Profiles

Until now, we looked at how to externalize application configuration to a property file, `application.properties`. What we want to be able to do is have different values for the same property in different environments.

Profiles provide a way to provide different configurations in different environments.

The following snippet shows how to configure an active profile in `application.properties`:

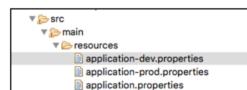
```
spring.profiles.active=dev
```

Once you have an active profile configured, you can define properties specific to that profile in `application-{profile-name}.properties`. For dev profile, the name of the properties file would be `application-dev.properties`. The following example shows the configuration in `application-dev.properties`:

```
application.enableSwitchForService1=true  
application.service1Url=http://abc-dev.service.com/somethingelse  
application.service1Timeout=250
```

The values in `application-dev.properties` will override the default configuration in `application.properties` if the active profile is dev.

We can have configurations for multiple environments, as shown here:



[ 239 ]

Copyrighted material

#### Advanced Spring Boot Features

### Profiles-based Bean configuration

Profiles can also be used to define different beans or different bean configurations in different environments. All classes marked with `@Component` or `@Configuration` can also be marked with an additional `@Profile` annotation to specify the profile in which the bean or configuration is enabled.

Let's consider an example. An application needs different caches enabled in different environments. In the dev environment, it uses a very simple cache. In production, we would want to use a distributed cache. This can be implemented using profiles.

The following bean shows the configuration enabled in a dev environment:

```
@Profile("dev")  
@Configuration  
public class DevSpecificConfiguration {  
    @Bean  
    public String cache() {  
        return "Dev Cache Configuration";  
    }  
}
```

The following bean shows the configuration enabled in a production environment:

```
@Profile("prod")  
@Configuration  
public class ProdSpecificConfiguration {  
    @Bean  
    public String cache() {  
        return "Production Cache Configuration - Distributed Cache";  
    }  
}
```

Based on the active profile configured, the respective configuration is picked up. Note that we are not really configuring a distributed cache in this example. We are returning a simple string to illustrate that profiles can be used to implement these kinds of variations.

### Other options for application configuration values

Until now, the approaches we took to configure application properties was using the key value pairs from either `application.properties` or `application-{profile-name}.properties`.

[ 240 ]

Copyrighted material

Pages 241 to 250 are not shown in this preview.

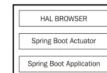
#### Advanced Spring Boot Features

### HAL Browser

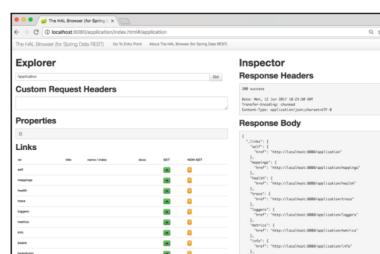
A number of these endpoints expose a lot of data. To be able to visualize the information better, we will add an HAL Browser to our application:

```
<dependency>  
    <groupId>org.springframework.data</groupId>  
    <artifactId>spring-data-rest-hal-browser</artifactId>  
</dependency>
```

Spring Boot Actuator exposes REST APIs around all the data captured from the Spring Boot application and environment. The HAL Browser enables visual representation around the Spring Boot Actuator API:



When we launch `http://localhost:8080/application` in the browser, we can see all the URLs exposed by actuator.



*Advanced Spring Boot Features*

Let's browse all the information exposed by actuator as part of different endpoints through the HAL Browser.

## Configuration properties

The configprops endpoint provides information about configuration options that can be configured through application properties. It basically is a collated list of all @ConfigurationProperties. The following screenshot shows configprops in HAL Browser:

```

{
  "endpoints": {
    "org.springframework.boot.actuate.endpoint.EndpointProperties": {
      "prefix": "spring",
      "properties": {
        "enabled": true
      }
    },
    "spring.transaction.org.springframework.boot.autoconfigure.transaction.TransactionProperties": {
      "prefix": "spring.transaction",
      "properties": {}
    },
    "management.info.org.springframework.boot.actuate.autoconfigure.InfoEndpointProperties": {
      "prefix": "management.info",
      "properties": {}
    }
  }
}
  
```

*Advanced Spring Boot Features*

To illustrate a known example, the following section from the service response shows the configuration options available for Spring MVC:

```

"spring.mvc": {
  "org.springframework.boot.autoconfigure.web.WebMvcProperties": {
    "prefix": "spring.mvc",
    "properties": {
      "dateFormat": null,
      "servlet": {
        "loadOnStartup": -1
      },
      "staticPathPattern": "**",
      "dispatchOptionsRequest": true,
      "dispatchTraceRequest": false,
      "locale": null,
      "ignoreDefaultModelOnRedirect": true,
      "logResolvedException": true,
      "async": {
        "requestTimeout": null
      },
      "messageCodesResolverFormat": null,
      "mediaTypes": {},
      "view": {
        "prefix": null,
        "suffix": null
      },
      "localeResolver": "ACCEPT_HEADER",
      "throwExceptionIfNoHandlerFound": false
    }
  }
}
  
```



To provide configuration for Spring MVC, we combine the prefix with the path in properties. For example, to configure loadOnStartup, we use a property with the name `spring.mvc.servlet.loadOnStartup`.

*Advanced Spring Boot Features*

## Environment details

The environment (env) endpoint provides information about the operating system, JVM installation, classpath, system environment variable, and the values configured in various application properties files. The following screenshot shows the environment endpoint in the HAL Browser:

```

{
  "os": {
    "name": "Linux"
  },
  "java": {
    "version": "1.8.0_111"
  },
  "process": {
    "pid": 1
  },
  "system": {
    "username": "root"
  }
}
  
```

**CUSTOM REQUEST HEADERS**

**Properties**

```
{
  "profiles": [
    "dev"
  ],
  "server": {
    "port": 8080
  },
  "systemProperties": {
    "java.awt.headless": "true",
    "java.awt.headless": "true"
  }
}
```

**Response Body**

```
{
  "profiles": [
    "dev"
  ],
  "server": {
    "port": 8080
  },
  "systemProperties": {
    "java.awt.headless": "true",
    "java.awt.headless": "true"
  }
}
```

[ 254 ]

Copyrighted material

Pages 255 to 266 are not shown in this preview.

### Spring Data

Some of the important Spring Data features are listed as follows:

- Easy integration with multiple data stores through various repositories
- The ability to parse and form queries based on repository method names
- Provides the default CRUD functionality
- Basic support for auditing, such as created by user and last changed by user
- Powerful integration with Spring
- Great integration with Spring MVC to expose REST controllers through **Spring Data Rest**

Spring Data is an umbrella project made up of a number of modules. A few of the important Spring Data modules are listed as follows:

- **Spring Data Commons:** Defines the common concepts for all Spring Data modules—repository and query methods
- **Spring Data JPA:** Provides easy integration with JPA repositories
- **Spring Data MongoDB:** Provides easy integration with MongoDB—a document-based data store
- **Spring Data REST:** Provides the functionality to expose Spring Data repositories as REST services with minimal code
- **Spring Data for Apache Cassandra:** Provides easy integration with Cassandra
- **Spring for Apache Hadoop:** Provides easy integration with Hadoop

In this chapter, we will take an in-depth look at the common concepts behind Spring Data, repository, and query methods. In the initial examples, we will use Spring Data JPA to illustrate these concepts. We will also take a look at a sample integration with MongoDB later in the chapter.

## Spring Data Commons

Spring Data Commons provides the basic abstractions behind Spring Data modules. We will use Spring Data JPA as an example to illustrate these abstractions.

Some of the important interfaces in Spring Data Commons are listed as follows:

```
Repository<T, ID extends Serializable>
CrudRepository<T, ID extends Serializable> extends Repository<T, ID>
PagingAndSortingRepository<T, ID extends Serializable> extends
CrudRepository<T, ID>
```

[ 267 ]

Copyrighted material

### Spring Data

## Repository

Repository is the core interface of Spring Data. It is a **marker interface**.

### The CrudRepository interface

The CrudRepository defines the basic Create, Read, Update, and Delete methods. The important methods in CrudRepository are shown in the following code:

```
public interface CrudRepository<T, ID extends Serializable>
extends Repository<T, ID> {
  <S extends T> S save(S entity);
  T findOne(ID primaryKey);
  Iterable<T> findAll();
  Long count();
  void delete(T entity);
  boolean exists(ID primaryKey);
  // ... more functionality omitted.
}
```

### The PagingAndSortingRepository interface

The PagingAndSortingRepository defines methods that provide the functionality to divide the ResultSet into pages as well as sort the results:

```
public interface PagingAndSortingRepository<T, ID extends
Serializable>
extends CrudRepository<T, ID> {
  Iterable<T> findAll(Sort sort);
  Page<T> findAll(Pageable pageable);
}
```

... [ 268 ] ...

## Spring Data JPA

Spring Data JPA implements the core functionality defined in Spring Data Common interfaces.

[ 268 ]

Copyrighted material

### Spring Data

`JpaRepository` is the JPA-specific repository interface:

```
public interface JpaRepository<T, ID extends Serializable>
    extends PagingAndSortingRepository<T, ID>, QueryByExampleExecutor<T> { }
```

`SimpleJpaRepository` is the default implementation of the `CrudRepository` interface for JPA:

```
public class SimpleJpaRepository<T, ID extends Serializable>
    implements JpaRepository<T, ID>, JpaSpecificationExecutor<T> { }
```

## Spring Data JPA example

Let's set up a simple project to understand the different concepts related to Spring Data Commons and Spring Data JPA.

The following are the steps involved:

1. Create a new project with `spring-boot-starter-data-jpa` as a dependency.
2. Add entities.
3. Add the `SpringBootApplication` class to run the application.
4. Create repositories.

### New project with Starter Data JPA

We will create a simple Spring Boot Maven project using the following dependencies:

```
<dependency>
    <groupId>org.springframework.boot</groupId>
    <artifactId>spring-boot-starter-data-jpa</artifactId>
</dependency>
<dependency>
    <groupId>com.h2database</groupId>
    <artifactId>h2</artifactId>
    <scope>runtime</scope>
</dependency>
<dependency>
    <groupId>org.springframework.boot</groupId>
    <artifactId>spring-boot-starter-test</artifactId>
    <scope>test</scope>
</dependency>
```

[ 269 ]

Copyrighted material

Pages 270 to 280 are not shown in this preview.

### Spring Data

We will start with examples related to finding rows matching specific attribute values. The following example shows different methods in order to search for the `User` by their name:

```
public interface UserRepository
    extends CrudRepository<User, Long> {
    List<User> findByName(String name);
    List<User> findByName(String name, Sort sort);
    List<User> findByName(String name, Pageable pageable);
    Long countByName(String name);
    Long deleteByName(String name);
    List<User> removeByName(String name);
}
```

Important things to note are as follows:

- `List<User> findByName(String name)`: The pattern is `findBy`, followed by the name of the attribute that you would want to query by. The value of the attribute is passed in as a parameter.
- `List<User> findByName(String name, Sort sort)`: This method allows you to specify a specific sort order.
- `List<User> findByName(String name, Pageable pageable)`: This method allows the use of pagination.
- Instead of `find` we can also use `read`, `query` or `get` to name the methods. For example, `queryByName` instead of `findByName`.
- Similar to `find..By` we can use `count..By` to find the count, and `delete..By` (or `remove..By`) to delete records.

The following example shows how to search by attributes of a containing element:

```
List<User> findByTodosTitle(String title);
```

The user contains `Todos`. `Todo` has `title` attribute. To create a method to search a user based on the title of the todo, we can create a method by the name `findByTodosTitle` in `UserRepository`.

The following examples show a few more variations that are possible with `findBy`:

```
public interface TodoRepository extends CrudRepository<Todo, Long> {
    List<Todo> findByTitleAndDescription
        (String title, String description);
    List<Todo> findDistinctTodoByTitleOrDescription
        (String title, String description);
    List<Todo> findByTitleIgnoreCase(String title, String
        description);
```

[ 281 ]

Copyrighted material

```
List<Todo> findByTitleOrderByIdDesc(String lastname);
List<Todo> findByIsDoneTrue(String lastname);
}
```

Important things to note are as follows:

- `findByTitleAndDescription`: Multiple attributes can be used to query
- `findDistinctTodoByTitleOrDescription`: Find distinct rows that match either title or description
- `findByTitleIgnoreCase`: Illustrates the use of the ignore case
- `findByTitleOrderByIdDesc`: Illustrates an example of specifying a specific sort order

The following example shows how to find a specific subset of records using find:

```
public interface UserRepository
    extends PagingAndSortingRepository<User, Long> {
    User findFirstByName(String name);
    User findTopByName(String name);
    List<User> findTopByName(String name);
    List<User> findFirst3ByName(String name);
}
```

Important things to note are as follows:

- `findFirstByName`, `findTopByName`: Queries for the first user
- `findTop3ByName`, `findFirst3ByName`: Finds the top three users

## Queries

Spring Data JPA also provides options to write custom queries. The following snippet shows a simple example:

```
@Query("select u from User u where u.name = ?1")
List<User> findUsersByNameUsingQuery(String name);
```

Important things to note are as follows:

- `@Query`: The annotation to define queries for repository methods
- `select u from User u where u.name = ?1`: Query to be executed. `?1` represents the first parameter
- `findUsersByNameUsingQuery`: When this method is called, the query specified is executed with the name as the parameter

— [ 282 ] —

Copyrighted material

Pages 283 to 284 are not shown in this preview.

Important things to note are as follows:

- `@RepositoryRestResource`: The annotation used to expose a repository using REST
- `collectionResourceRel = "users"`: The `collectionResourceRel` value to be used in the generated links
- `path = "users"`: The path under which the resource has to be exposed

When we launch `SpringDataJpaFirstExampleApplication` as a Java application, the following can be seen in the log:

```
s.b.c.e.t.TomcatEmbeddedServletContainer : Tomcat initialized with port(s):
    8080 (http)
o.s.b.w.servlet.ServletRegistrationBean : Mapping servlet:
    'dispatcherServlet' to {}
o.s.b.w.servlet.FilterRegistrationBean : Mapping filter:
    'characterEncodingFilter' to: [/]
o.s.w.m.m.a.RequestMappingHandlerMapping : Mapped "[/{repository}]" onto ***
o.s.d.r.w.RepositoryRestHandlerMapping : Mapped "[/{repository}]",
methods=[OPTIONS]
o.s.d.r.w.RepositoryRestHandlerMapping : Mapped "[/{repository}]",
methods=[HEAD]
o.s.d.r.w.RepositoryRestHandlerMapping : Mapped "[/{repository}]",
methods=[GET]
o.s.d.r.w.RepositoryRestHandlerMapping : Mapped "[/{repository}]",
methods=[POST]
o.s.d.r.w.RepositoryRestHandlerMapping : Mapped "[/{repository}/{id}]",
methods=[OPTIONS]
o.s.d.r.w.RepositoryRestHandlerMapping : Mapped "[/{repository}/{id}/{property}]"
o.s.d.r.w.RepositoryRestHandlerMapping : Mapped "[/{repository}/{search}]",
methods=[GET]
```

The preceding log shows that the Spring MVC DispatcherServlet is launched and ready to serve different request methods and URLs.

— [ 285 ] —

Copyrighted material

Pages 286 to 299 are not shown in this preview.

Important things to note are as follows:

- `@Value("${number.service.url}") private String numberServiceUrl`: We would want the number service URL to be configurable in application properties.
- `@RequestMapping("/add") public Long add()`: Exposes a service at the URI /add. The add method calls the number service using `RestTemplate` and has the logic to sum the numbers returned in the response.

Let's configure `application.properties`, as shown in the following snippet:

```
spring.application.name=service-consumer
server.port=8100
```

Important things to note are as follows:

- `spring.application.name=service-consumer`: Configures a name for the Spring Boot application
- `server.port=8100`: Uses 8100 as the port for service consumer
- `number.service.url=http://localhost:8080/random`: Configures the number service URL for use in the add service

When the service is called at the URL `http://localhost:8100/add`, the following response is returned:

2890

The following is an extract from the log of Microservice A:

```
c.m.s.c.c.RandomNumberController : Returning [752, 119, 493, 871, 445]
```

The log shows that random service from Microservice A returned 5 numbers. The add service in the service consumer added them up and returned a result 2890.

We now have our example microservices ready. In the next steps, we will add Cloud-Native features to these microservices.

## Ports

In this chapter, we will create six different microservices applications and components. To keep things simple, we will use specific ports for specific applications.

[ 300 ]

Copyrighted material

Pages 301 to 330 are not shown in this preview.

## Spring Cloud

- `public boolean shouldFilter()`: If the filter should only be executed in certain conditions, the logic can be implemented here. If you would want the filter to always be executed, return `true`.
- `public Object run()`: The method to implement the logic for the filter. In our example, we are logging the request method and the URL of the request.

When we start up the Zuul server by launching `ZuulApiGatewayServerApplication` as a Java application, you will see the following log in Eureka Name Server:

```
Registered instance ZUUL-API-GATEWAY/192.168.1.5:zuul-api-gateway:8765 with status UP (replication=false)
```

This shows that Zuul API gateway is up and running. Zuul API gateway is also registered with Eureka Server. This allows microservice consumers to talk to the Name server to get details about Zuul API gateway.

The following figure shows the Eureka dashboard at `http://localhost:8761`. You can see that instances of Microservice A, service consumer, and Zuul API Gateway are now registered with Eureka Server:

Instances currently registered with Eureka			
Application	Address	Availability Zones	Status
MICROSERVICE-A	n/a (1)	(1)	UP (1) - 192.168.1.5:microservice-a
SERVICE-CONSUMER	n/a (1)	(1)	UP (1) - 192.168.1.5:service-consumer:8100
ZUUL-API-GATEWAY	n/a (1)	(1)	UP (1) - 192.168.1.5:zuul-api-gateway:8765

The following is an extract from the Zuul API gateway log:

```
Mapped URL path [/microservice-a/*] onto handler of type [class org.springframework.cloud.netflix.zuul.web.ZuulController]
Mapped URL path [/service-consumer/*] onto handler of type [class org.springframework.cloud.netflix.zuul.web.ZuulController]
```

By default, all services in Microservice A and the service consumer microservice are enabled for reverse proxying by Zuul.

[ 331 ]

Copyrighted material

## Spring Cloud

### Invoking microservices through Zuul

Let's invoke random service through the service proxy now. The direct URL to a random microservice is `http://localhost:8080/random`. This is exposed by Microservice A, whose application name is `microservice-a`.

The URL structure to call a service through Zuul API Gateway is `http://localhost:(port)/{microservice-application-name}/{service-uri}`. So, the Zuul API Gateway URL for random service is `http://localhost:8765/microservice-a/random`. When you invoke random service through API Gateway, you get a response shown here. The response is similar to what you would typically get when directly calling the random service:

```
[73, 671, 339, 354, 211]
```

The following is an extract from the Zuul API Gateway log. You can see that the SimpleLoggingFilter that we created in Zuul API Gateway is executed for the request:

```
c.m.s.z.filters.pre.SimpleLoggingFilter : Request Method : GET
URL: http://localhost:8765/microservice-a/random
```

The add service is exposed by service consumer, whose application name is service-consumer and the service URI is /add. So, the URL to execute the add service through the API Gateway is `http://localhost:8765/service-consumer/add`. The response from the service is shown here. The response is similar to what you would typically get when directly calling the add service:

2488

The following is an extract from the Zuul API Gateway log. You can see that the initial add service call is going through the API Gateway:

```
2017-03-28 14:05:17.514 INFO 83147 --- [nio-8765-exec-1]
c.m.s.z.filters.pre.SimpleLoggingFilter : Request Method : GET
URL: http://localhost:8765/service-consumer/add
```

The add service calls random service on Microservice A. While the initial call to add service goes through the API Gateway, the call from add service (service consumer microservice) to random service (Microservice A) is not routed through API Gateway. In an ideal world, we would want all the communication to take place through API Gateway.

[ 332 ]

Copyrighted material

Pages 333 to 342 are not shown in this preview.

### Spring Cloud

NumberAdderController exposes a service with request mapping /add. This uses RandomServiceProxy to fetch random numbers. What if this service fails? Hystrix provides a fallback. The following snippet shows how we can add a fallback method to a request mapping. All we need to do is add the @HystrixCommand annotation to the fallbackMethod attribute, defining the name of the fallback method--in this example, getDefaultResponse:

```
@HystrixCommand(fallbackMethod = "getDefaultResponse")
@RequestMapping("/add")
public Long add() {
    //Logic of add() method
}
```

Next, we define the getDefaultResponse() method with the same return type as the add() method. It returns a default hardcoded value:

```
public Long getDefaultResponse() {
    return 10000L;
}
```

Let's bring down Microservice A and invoke <http://localhost:8100/add>. You will get the following response:

10000

When Microservice A fails, the service consumer microservice handles it gracefully and offers reduced functionality.

## Summary

Spring Cloud makes it easy to add Cloud-Native features to your microservices. In this chapter, we looked at some of the important patterns in developing Cloud-Native applications and implemented them using various Spring Cloud projects.

It is important to remember that the field of developing Cloud-Native applications is still in its inception phase--in its first few years. It would need more time to mature. Expect some evolution in patterns and frameworks in the years to come.

In the next chapter, we will shift our attention to Spring Data Flow. Typical use cases on the cloud include real-time data analytics and data pipelines. These use cases involve the flow of data between multiple microservices. Spring Data Flow provides patterns and best practices for distributed streaming and data pipelines.

[ 343 ]

Copyrighted material

# 10

## Spring Cloud Data Flow

Spring Data Flow brings the microservices architecture into typical data flow and event flow scenarios. We will discuss more about these scenarios later in this chapter. Building on top of other Spring Projects, such as Spring Cloud Stream, Spring Integration, and Spring Boot, Spring Data Flow makes it easy to define and scale use cases involving data and event flows using message-based integration.

In this chapter, we will discuss the following topics:

- Why do we need asynchronous communication?
- What is Spring Cloud Stream? How does it build on top of Spring Integration?
- Why do we need Spring Data Flow?
- What are the important concepts in Spring Data Flow you would need to understand?
- What are the use cases where Spring Data Flow is useful?

We will also implement a simple event flow scenario with three microservices acting as the source (application generating the events), processor, and sink (application consuming events). We will implement the microservices using Spring Cloud Stream and establish connections between them over the message broker using Spring Cloud Data Flow.

### Message-based asynchronous communication

There are two options when integrating applications:

- **Synchronous:** Service consumer invokes the service provider and waits for a response.

Copyrighted material

Pages 345 to 356 are not shown in this preview.

Update the `SpringBootApplication` file with the following code:

```

@EnableBinding(Processor.class)
public class StockIntelligenceProcessorApplication {
    private static final Logger logger = LoggerFactory.getLogger(StockIntelligenceProcessorApplication.class);
    private static Map<StockTicker, Integer> holdings = getHoldingsFromDatabase();
    private static Map<StockTicker, Integer> getHoldingsFromDatabase() {
        final Map<StockTicker, Integer> holdings = new HashMap<>();
        holdings.put(StockTicker.FACEBOOK, 10);
        holdings.put(StockTicker.GOOGLE, 0);
        holdings.put(StockTicker.IBM, 15);
        holdings.put(StockTicker.MICROSOFT, 30);
        holdings.put(StockTicker.TWITTER, 50);
        return holdings;
    }
    @Transformer(inputChannel = Processor.INPUT,
    outputChannel = Processor.OUTPUT)
    public Object addOurInventory(StockPriceChangeEvent event) {
        logger.info("Started processing event " + event);
        Integer holding = holdings.get(StockTicker.valueOf(event.getStockTicker()));
        StockPriceChangeEventWithHoldings eventWithHoldings = new StockPriceChangeEventWithHoldings(event, holding);
        logger.info("Ended processing eventWithHoldings " +
        + eventWithHoldings);
        return eventWithHoldings;
    }
    public static void main(String[] args) {
        SpringApplication.run(
            StockIntelligenceProcessorApplication.class, args);
    }
}

```

A few important things to note are as follows:

- `@EnableBinding(Processor.class)`: The `EnableBinding` annotation enables binding a class with the respective channel it needs—an input and/or an output. The `Processor` class is used to register a Cloud Stream with one input channel and one output channel.

[ 357 ]

Copyrighted material

Pages 358 to 369 are not shown in this preview.

You should see the following output if the stream is successfully created:

Created new stream 'process-stock-change-events'

You can also see the registered stream on the **Streams** tab of Spring Cloud Data Flow dashboard at <http://localhost:9393/dashboard>, as shown in the following screenshot:

The screenshot shows the 'Streams' tab of the Spring Cloud Data Flow dashboard. At the top, there are tabs for APPS, RUNTIME, STREAMS, TASKS, JOBS, ANALYTICS, and ABOUT. The STREAMS tab is selected. Below the tabs, there's a search bar labeled 'Search definitions'. Underneath, there are two buttons: 'Definitions' (selected) and 'Create Stream'. A table lists the stream definition: 'Name' is 'process-stock-change-events', 'Definition' is 'significane-stock-change-source->stock-intelligence->process-event-store-sink', 'Status' is 'undeployed', and 'Actions' include 'Details', 'Undeploy', 'Deploy', and 'Destroy'. There are also 'Expand All' and 'Collapse All' buttons.

## Deploying the stream

To deploy the stream, we can execute the following command on the Data Flow Shell:

```
stream deploy --name process-stock-change-events
```

You will see the message shown here when the request is sent for the creation of stream:

Deployment request has been sent for stream 'process-stock-change-events'

The following extract shows an extract from the Local Data Flow Server log:

```

o.s.c.d.spi.local.LocalAppDeployer : deploying app process-stock-change-
events.event-store-sink instance 0
Logs will be in /var/folders/y/_x4jdvd7w94q5gsh745gzz0000gn/T/spring-
cloud-dataflow-3084432375250471462/process-stock-change-
events-1492100265496/process-stock-change-events.event-store-sink
o.s.c.d.spi.local.LocalAppDeployer : deploying app process-stock-change-
events.stock-intelligence-processor instance 0

```

[ 370 ]

Copyrighted material

Pages 371 to 465 are not shown in this preview.